

From Chef to Docker

A Saner Management for a Multi-facet Development Environment

Whoami

Yoav Landman
JFrog CTO



What Frog?



What Frog?

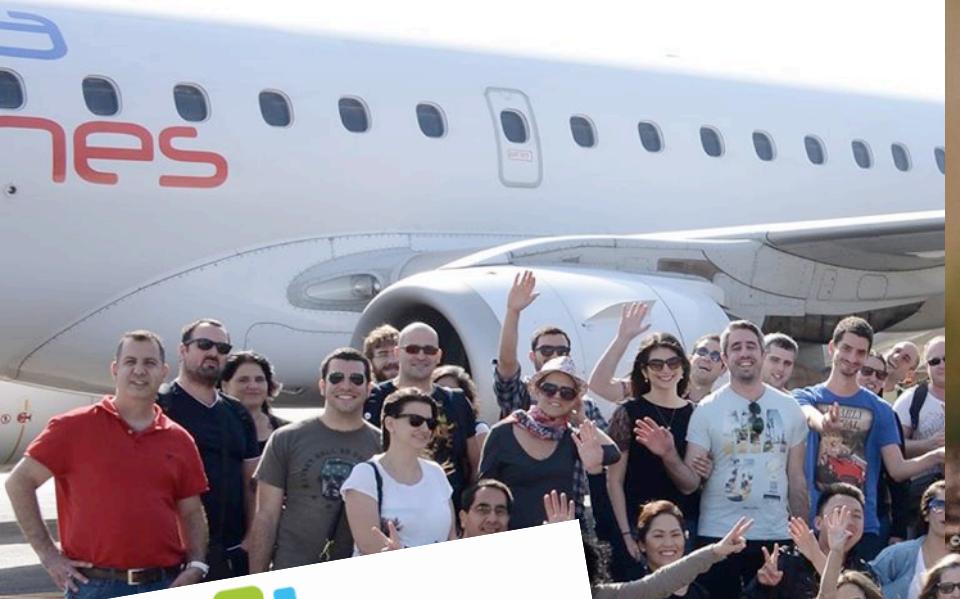


What Frog?



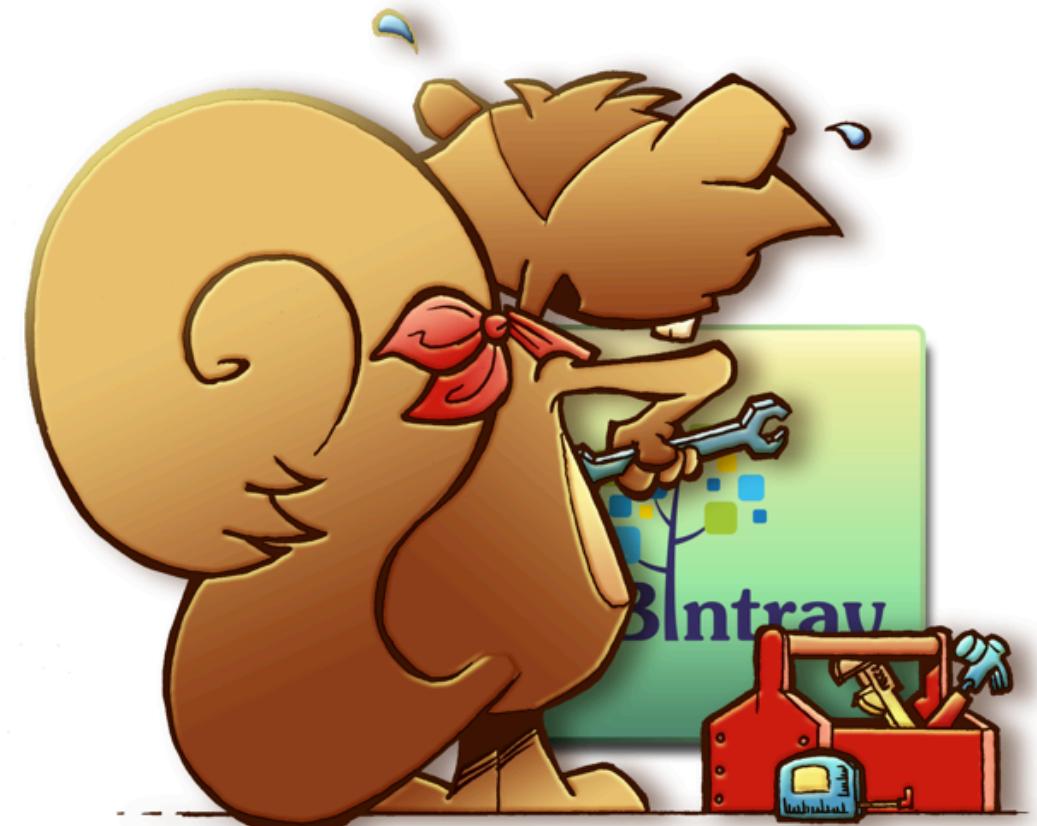
artifactory

What Frog?



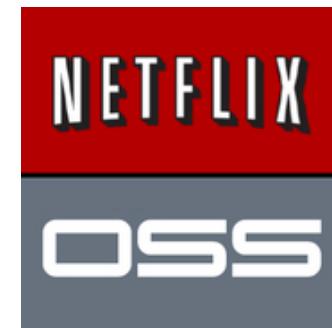


So...



Bintray

Distribution as a service



ThoughtWorks®



Remember 2013

Developing for multi-component environments while keeping your sanity

A session at DayNexus 2013



Fred Simon

Monday 18th February, 2013 2:30pm to 3:45pm (EST)

Managing a modern multi-component application in a continuous integration/deployment environment can be very tricky.

In this session we will concentrate on two of the aspects - environment setup, showcasing tools like Vagrant, Chef and Puppet for creating and controlling development, testing, staging, production and other environments and pushing the deliverables through those environments, using tools like Gradle, Jenkins and Artifactory.

Join us and get all the little tips and tricks that will help you rule your development world.

About the speaker



Fred Simon

JFrog and AlphaCSP founder. [bio from Twitter](#)

Untrack

Don't meet

Coverage of this session



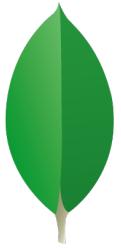
SLIDES

[Developing for multi-component environments while keeping your sanity \(speakerdeck.com\)](#)

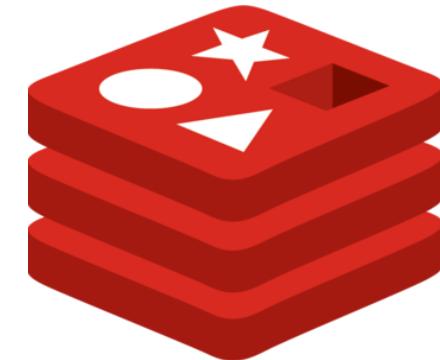
★ Favourite this? Added 2 years ago edit delete



Diverse technologies

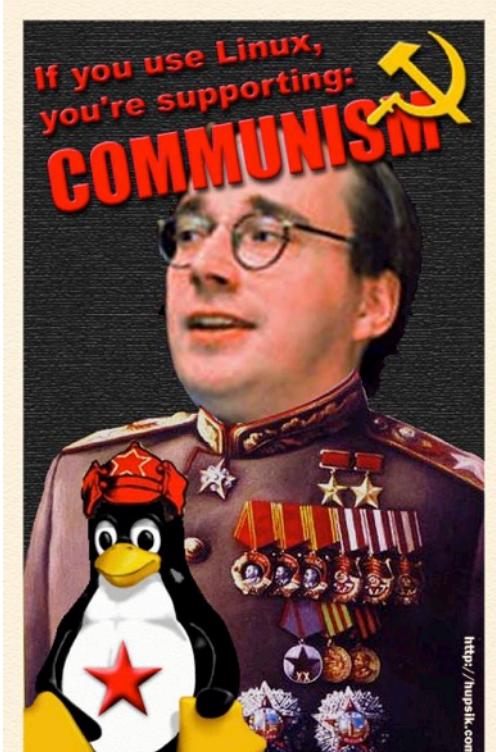


mongoDB

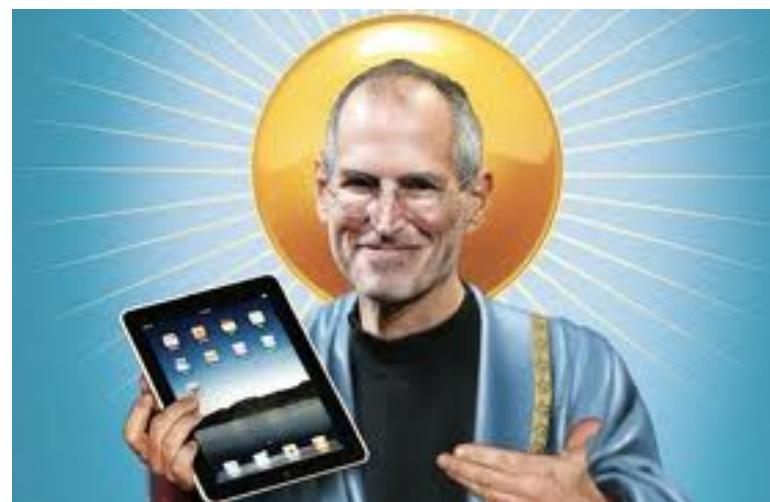
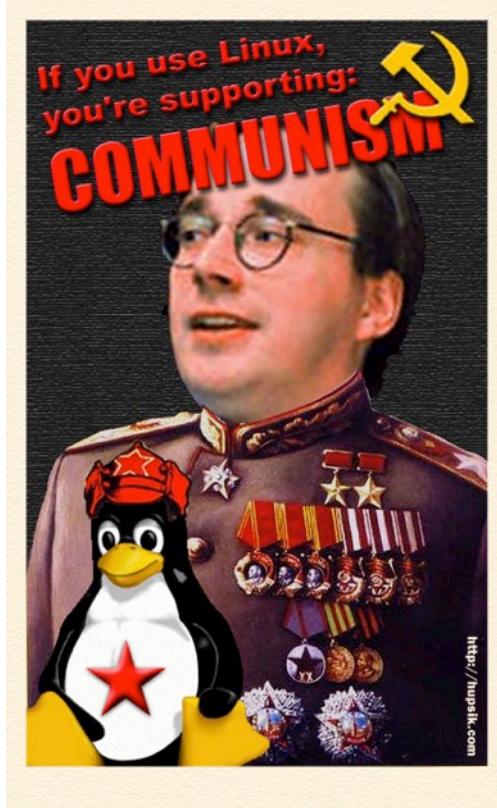


NGINX

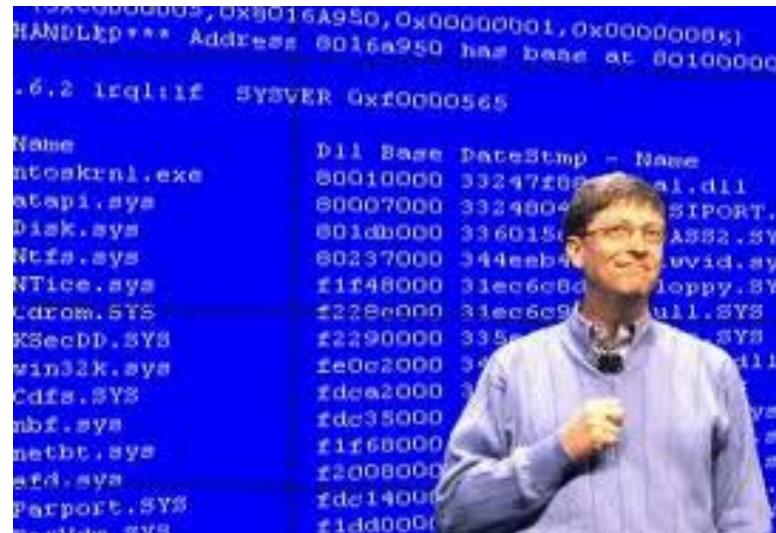
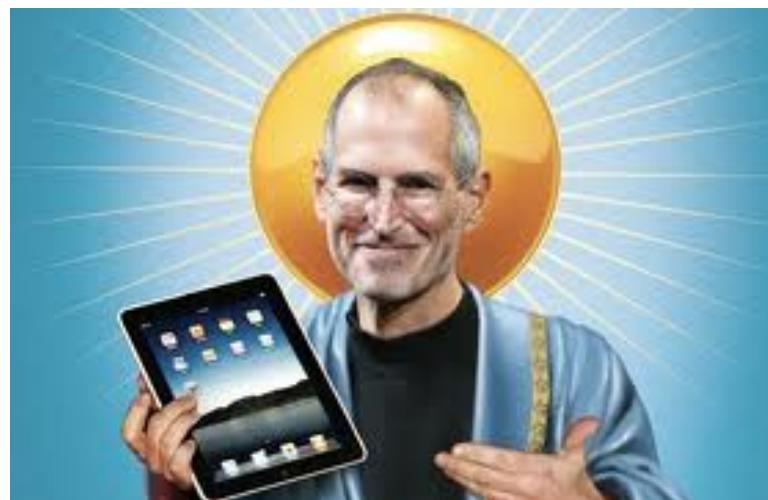
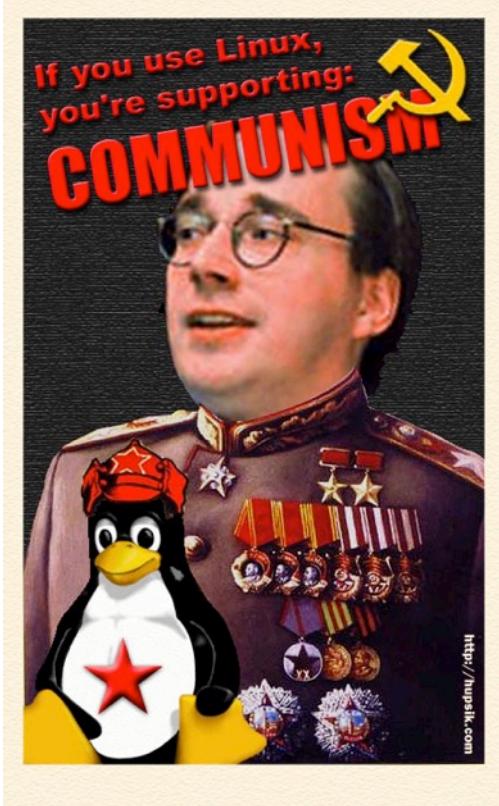
We are liberal



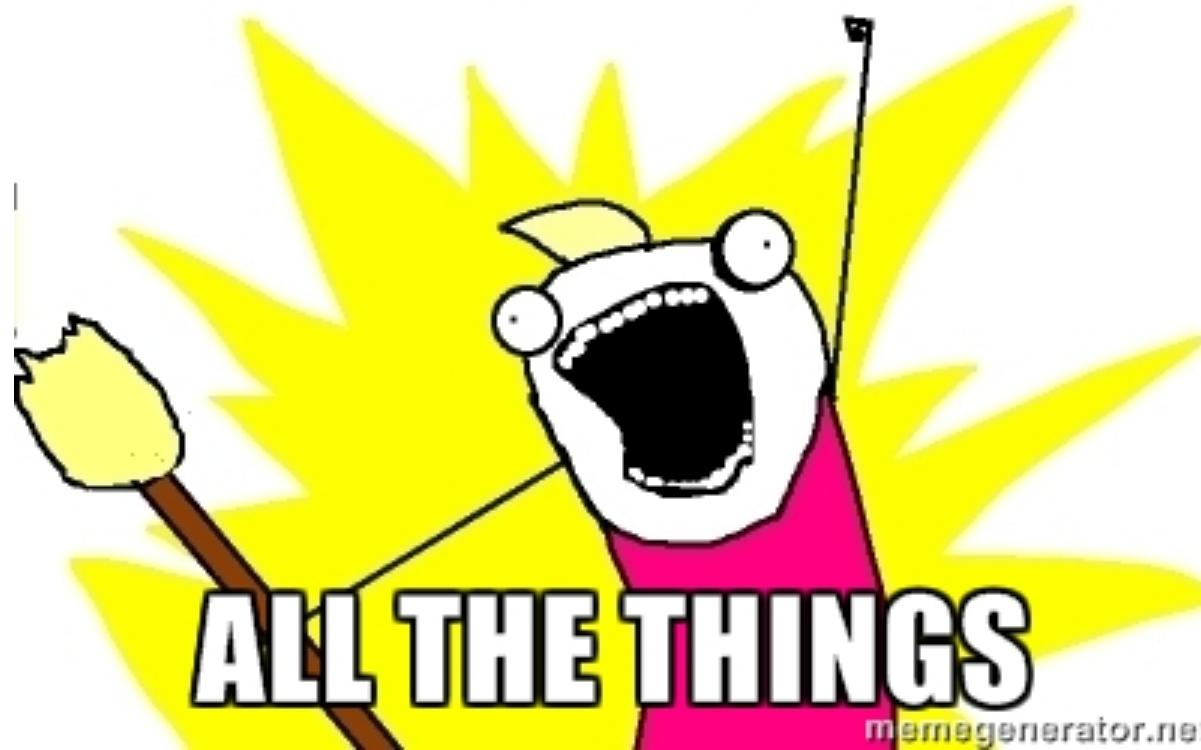
We are liberal



We are liberal

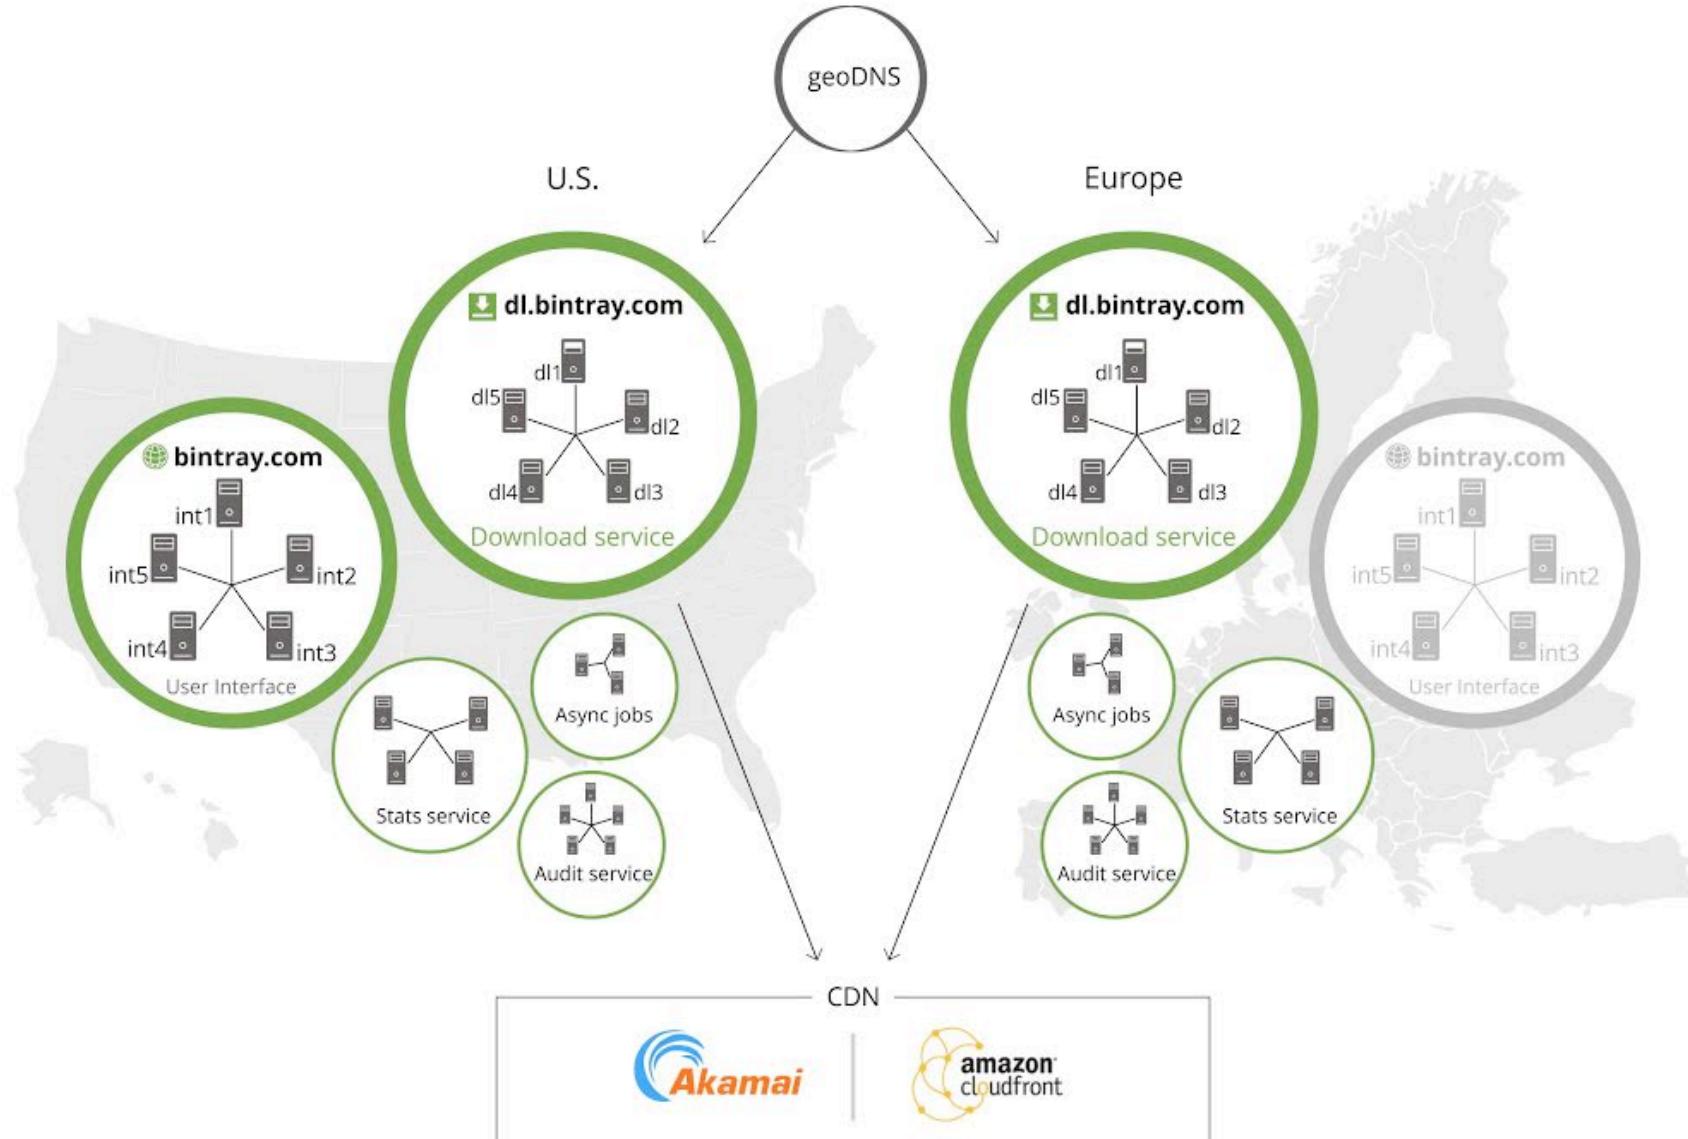


MICROSERVICE



ALL THE THINGS

memegenerator.net



A close-up photograph of Aragorn's face from the Lord of the Rings movies. He has long, dark hair and a beard, looking slightly to the side with a thoughtful expression. The background is blurred.

ONE DOES NOT SIMPLY

SETUP ALL THAT ON ALL THIS

OK, we have an idea

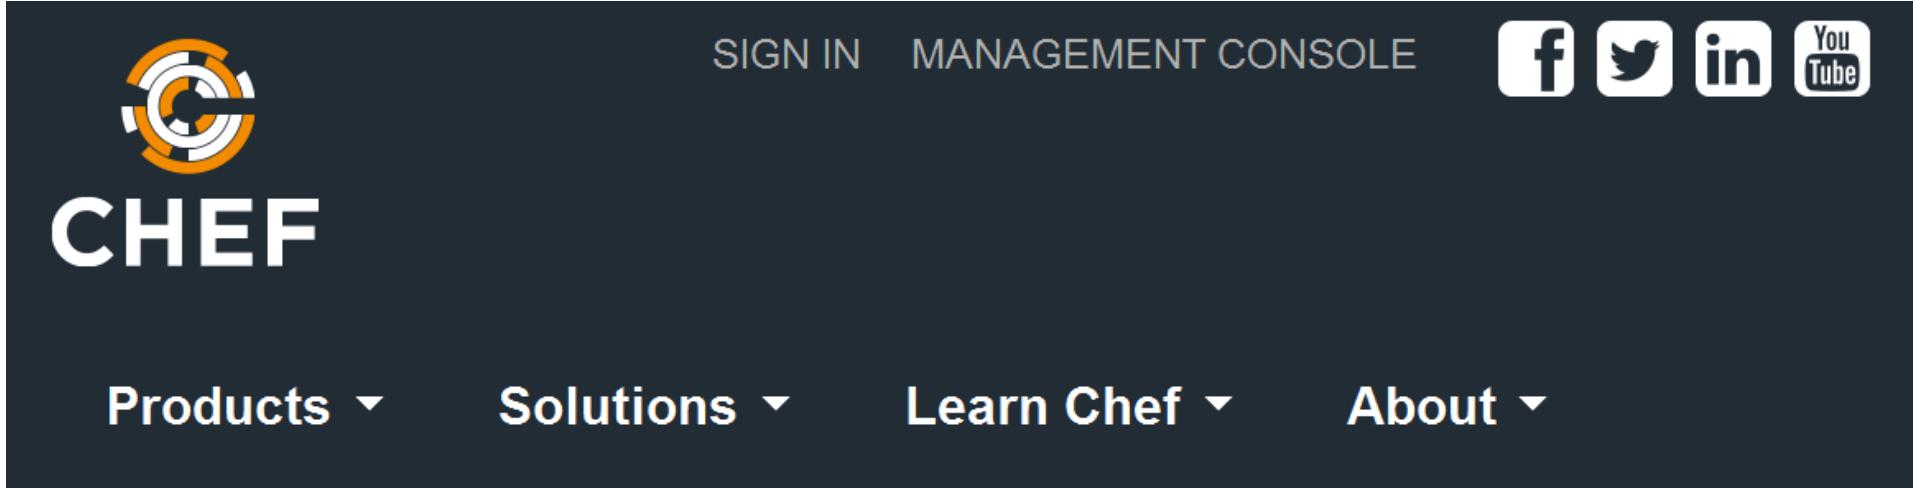


The Solution



CHEF™

Opscode Chef



The image shows the top navigation bar of the Opscode Chef website. It features the Chef logo (a stylized orange and white gear icon) and the word "CHEF" in large white capital letters on the left. On the right, there are links for "SIGN IN" and "MANAGEMENT CONSOLE". Below these are social media icons for Facebook, Twitter, LinkedIn, and YouTube. A horizontal menu bar below the navigation bar contains four dropdown items: "Products", "Solutions", "Learn Chef", and "About".

Chef is an automation platform that transforms infrastructure into code. Stop thinking in terms of physical and virtual servers. With Chef, your real asset is the code that brings those servers and the services they provide to life. An automated infrastructure can accelerate your time to

Chef works with cookbooks

Cookbooks

A cookbook is the fundamental **unit of configuration** and policy distribution.

Each cookbook defines a **scenario**, such as everything needed to install and configure MySQL, and then it contains all of the components that are required to support that scenario, including:

- Attribute values that are set on nodes
- Definitions that allow the creation of reusable collections of resources
- File distributions

Chef works with cookbooks

Cookbooks

A cookbook is the fundamental unit of configuration and policy distribution. Each cookbook defines a scenario, such as everything needed to install and configure MySQL, and then it contains all of the components that are required to support that scenario, including:

- Attribute values that are set on nodes
- Definitions that allow the creation of reusable collections of resources
- File distributions

Example: Install MySQL in Ruby

```
115
116  if platform_family? 'windows'
117    require 'win32/service'
118
119    windows_path node['mysql']['bin_dir'] do
120      action :add
121    end
122
123    windows_batch "install mysql service" do
124      command "\"#{node['mysql']['bin_dir']}\\mysqld.exe\" --install #{node['mysql']['service_name']}"
125      not_if { Win32::Service.exists?(node['mysql']['service_name']) }
126    end
127  end
128
```

Vagrant



VAGRANT

HOW VAGRANT BENEFITS YOU

If you're a **developer**, Vagrant will isolate dependencies and their configuration within a single disposable, consistent environment, without sacrificing any of the tools you're used to working with (editors, browsers, debuggers, etc.). Once you or someone else creates a single Vagrantfile, you just need to `vagrant up` and everything is installed

Setup Vagrant with Chef

CHEF SOLO PROVISIONER

Provisioner name: `chef_solo`

The chef solo provisioner allows you to provision the guest using [Chef](#), specifically with [Chef Solo](#).

Setup Vagrant with Chef

```
config.vm.box = "centos-6"

# Enable and configure the chef solo provisioner
config.vm.provision :chef_solo do |chef|
  chef.cookbooks_path = "#{checkout}"
  #chef.data_bags_path = "install/chef/data_bags"
  chef.roles_path = "install/chef/roles"

  chef.add_role("bintray_common_vagrant_server")
  chef.add_role("bintray_interaction_vagrant_server")
  chef.add_role("bintray_download_vagrant_server")
end
end
```

Setup Vagrant with Chef

```
config.vm.box = "centos-6"
```

```
# Enable and configure the chef solo provisioner
config.vm.provision :chef_solo do |chef|
  chef.cookbooks_path = "#{checkout}"
  #chef.data_bags_path = "install/chef/data_bags"
  chef.roles_path = "install/chef/roles"

  chef.add_role("bintray_common_vagrant_server")
  chef.add_role("bintray_interaction_vagrant_server")
  chef.add_role("bintray_download_vagrant_server")
end
end
```

Setup Vagrant with Chef

```
config.vm.box = "centos-6"

# Enable and configure the chef solo provisioner
config.vm.provision :chef_solo do |chef|
  chef.cookbooks_path = "#{checkout}"
  #chef.data_bags_path = "install/chef/data_bags"
  chef.roles_path = "install/chef/roles"

  chef.add_role("bintray_common_vagrant_server")
  chef.add_role("bintray_interaction_vagrant_server")
  chef.add_role("bintray_download_vagrant_server")
end
end
```

Setup Vagrant with Chef

```
config.vm.box = "centos-6"

# Enable and configure the chef solo provisioner
config.vm.provision :chef_solo do |chef|
  chef.cookbooks_path = "#{checkout}"
  #chef.data_bags_path = "install/chef/data_bags"
  chef.roles_path = "install/chef/roles"

  chef.add_role("bintray_common_vagrant_server")
  chef.add_role("bintray_interaction_vagrant_server")
  chef.add_role("bintray_download_vagrant_server")
end
end
```

Setup Vagrant with Chef

```
config.vm.box = "centos-6"

# Enable and configure the chef solo provisioner
config.vm.provision :chef_solo do |chef|
  chef.cookbooks_path = "#{checkout}"
  #chef.data_bags_path = "install/chef/data_bags"
  chef.roles_path = "install/chef/roles"

  chef.add_role("bintray_common_vagrant_server")
  chef.add_role("bintray_interaction_vagrant_server")
  chef.add_role("bintray_download_vagrant_server")
end
end
```

How it works

How it works

1. Check out cookbook and vagrant file

How it works

1. Check out cookbook and vagrant file
2. Vagrant downloads box

How it works

1. Check out cookbook and vagrant file
2. Vagrant downloads box
3. Vagrant starts the box

How it works

1. Check out cookbook and vagrant file
2. Vagrant downloads box
3. Vagrant starts the box
4. Vagrant triggers Chef

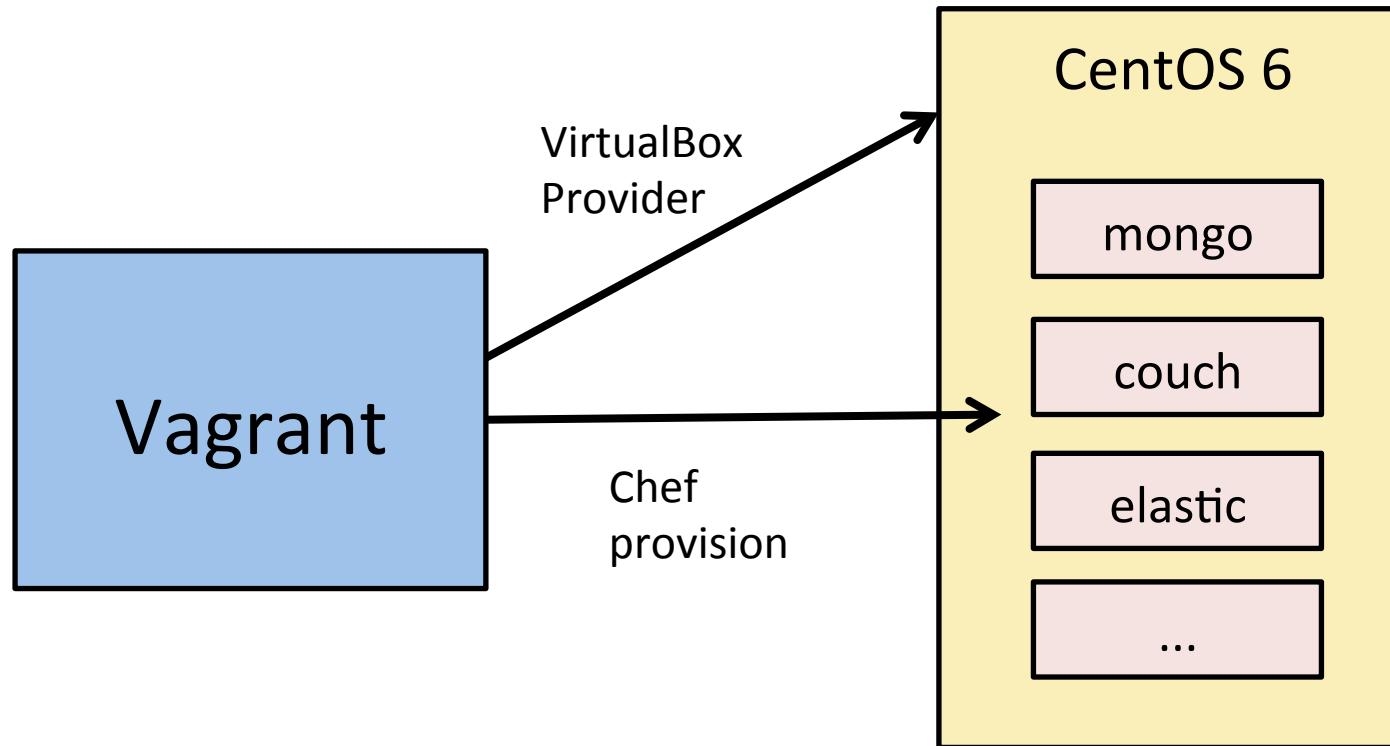
How it works

1. Check out cookbook and vagrant file
2. Vagrant downloads box
3. Vagrant starts the box
4. Vagrant triggers Chef
5. Chef plays recipes to set up everything

How did we like it?



Solves the problem!



But...



Developers break things

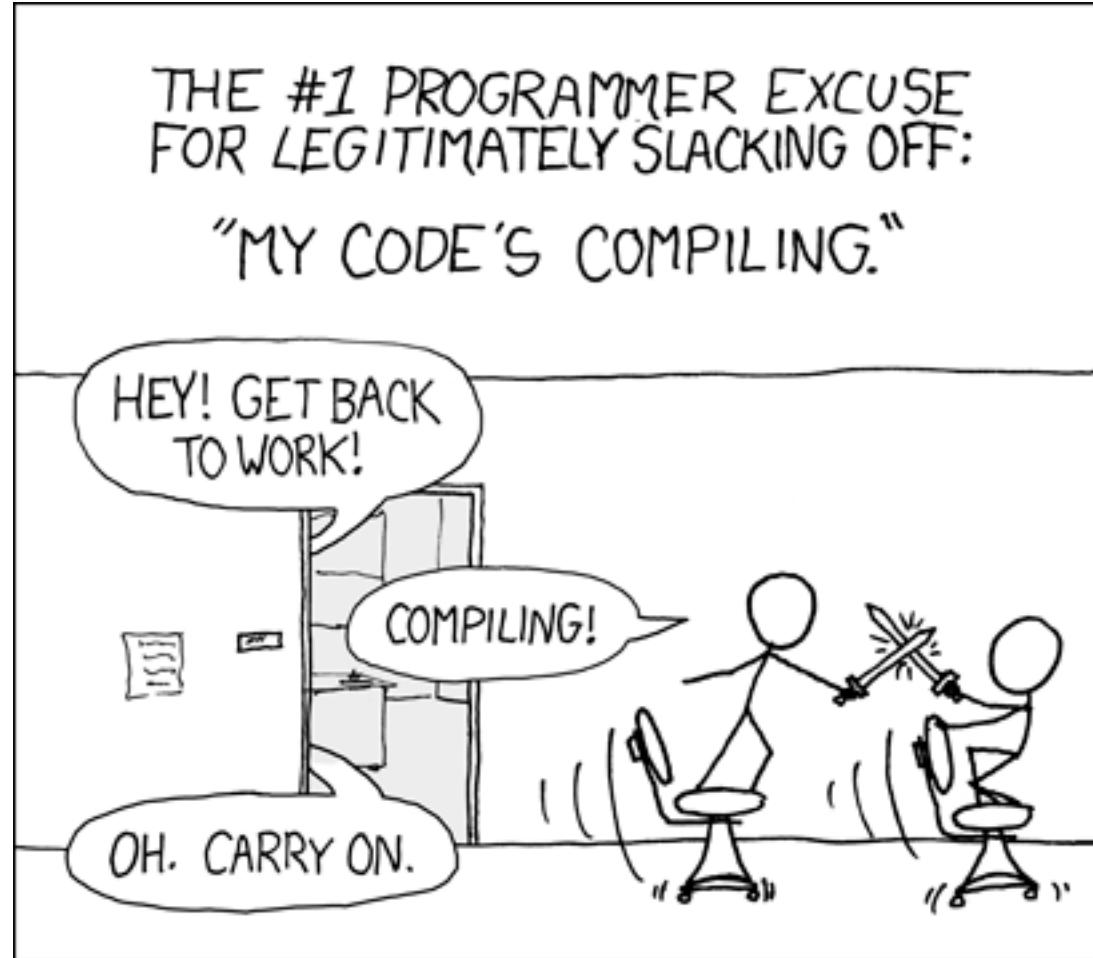




ONE DOES NOT SIMPLY

REPLAY ALL CHEF RECIPES

I am not slacking off, my environment is provisioned



Solution I:

Pre-build instead of
build

Not really new idea



How it works

How it works

1. Vagrant downloads box

How it works

1. Vagrant downloads box
2. Vagrant starts the box

How it works

1. Vagrant downloads box
2. Vagrant starts the box
3. Everything is up and running!

How did we like it?



Solves the problem!

But every time something changing
in the environment...



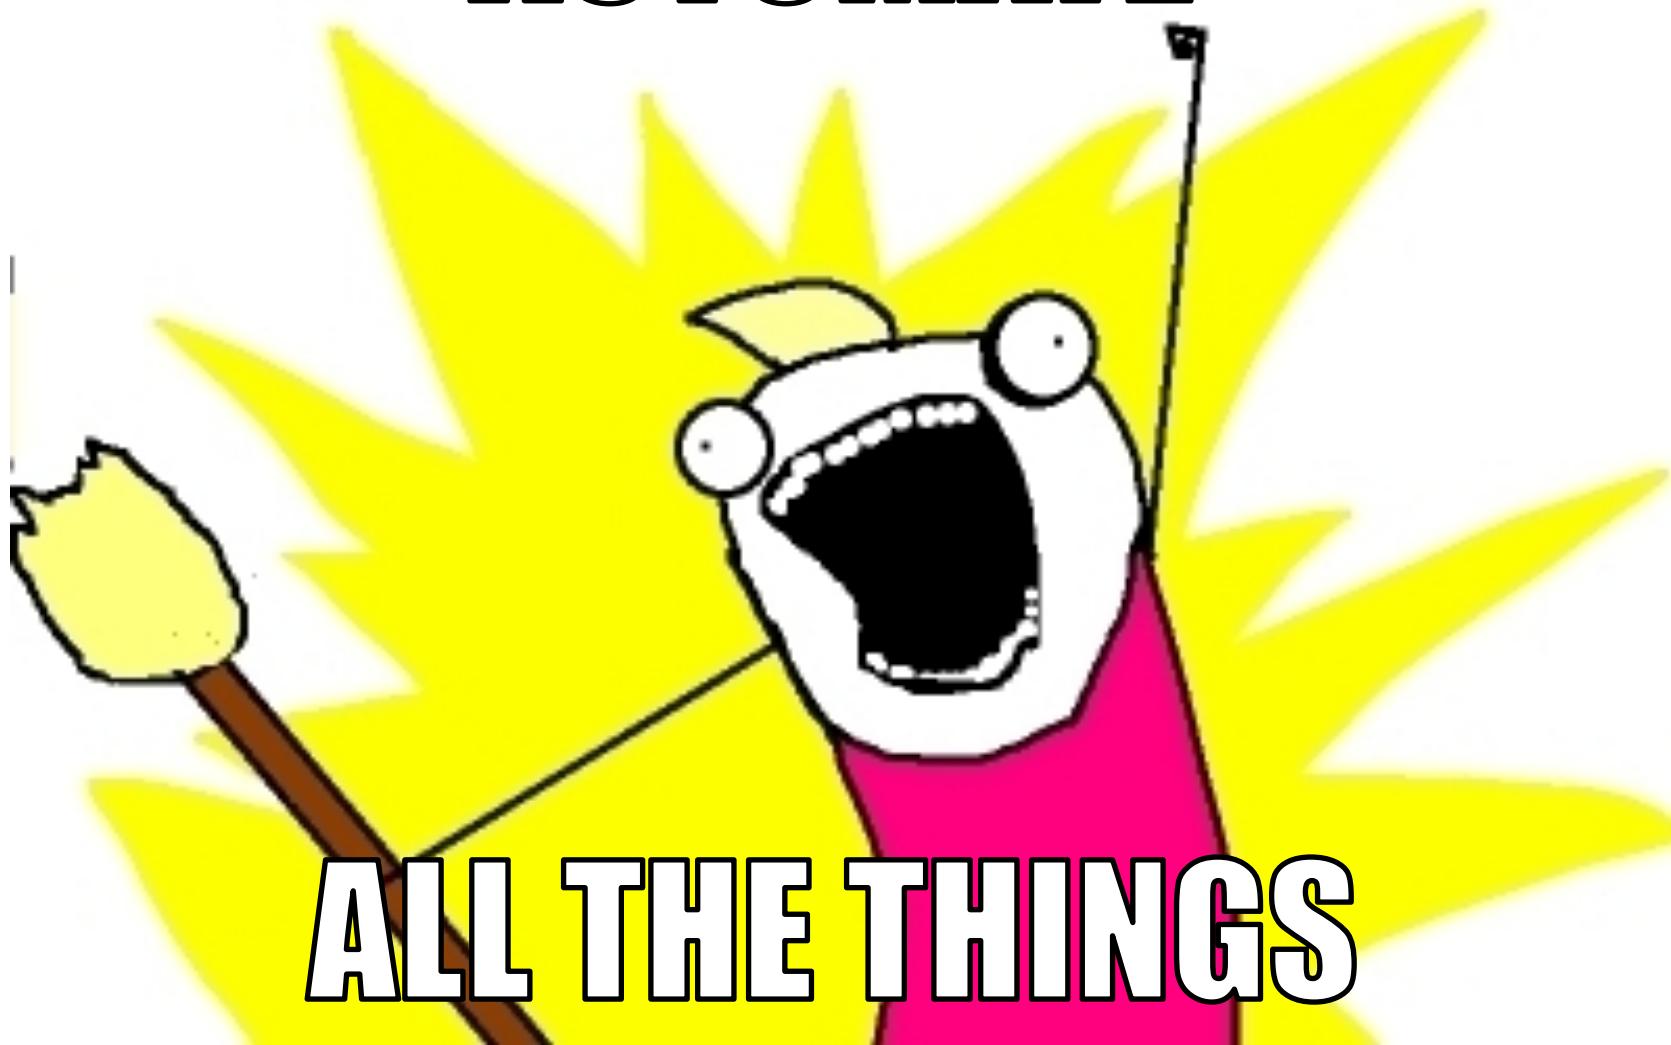


**ONE DOES NOT SIMPLY
REBUILD VM ON EVERY CHANGE**

OK, we have an idea



AUTOMATE



ALL THE THINGS

Automating vm creation

Automating vm creation

Chef?

Automating vm creation

Packer?

Automating vm creation

docker?

A scene from Toy Story featuring Woody and Buzz Lightyear. Woody, on the left, has a neutral, slightly weary expression. Buzz, on the right, is in mid-air, performing a handstand or trick, with his arms raised and legs bent. He is wearing his signature green space ranger suit with purple stripes and a red button on the chest. The background shows a room with a chalkboard and some toys.

DOCKER

DOCKER EVERYWHERE

WHY DOCKER?



Operating-system-level virtualization

From Wikipedia, the free encyclopedia

(Redirected from [Software container](#))

Operating-system-level virtualization is a server [virtualization](#) method where the [kernel](#) of an [operating system](#) allows for multiple isolated [user space](#) instances, instead of just one. Such instances (often called [containers](#), [virtualization engines](#) (VE), [virtual private servers](#) (VPS), or [jails](#)) may look and feel like a real server from the point of view of its owners and users.

On [Unix-like](#) operating systems, this technology can be seen as an advanced implementation of the standard [chroot](#) mechanism. In addition to isolation mechanisms, the kernel often provides resource management features to limit the impact of one container's activities on the other containers.

Contents [\[show\]](#)

Uses [\[edit\]](#)

Operating-system-level virtualization is commonly used in [virtual hosting](#) environments, where it is useful for securely allocating finite hardware resources amongst a large number of mutually-distrusting users. System administrators may also use it, to a lesser extent, for consolidating server hardware by moving services on separate hosts into containers on the one server.

Other typical scenarios include separating several applications to separate containers for improved security, hardware independence, and added resource management features. The improved security provided by the use of a chroot mechanism, however, is nowhere near ironclad.^[1] Operating-system-level virtualization implementations capable of [live migration](#) can also be used for dynamic load balancing of containers between nodes in a cluster.

Overhead [\[edit\]](#)

Operating-system-level virtualization usually imposes little to no overhead, because programs in virtual partitions use the operating system's normal [system call](#) interface and do not need to be subjected to [emulation](#) or be run in an intermediate [virtual machine](#), as is the case with whole-system virtualizers (such as [VMware ESXi](#), [QEMU](#) or [Hyper-V](#)) and paravirtualizers (such as [Xen](#) or [UML](#)). This form of virtualization also does not require support in hardware to perform efficiently.

Operating-system-level virtualization

From Wikipedia, the free encyclopedia

(Redirected from [Software container](#))

Operating-system-level virtualization is a server virtualization method where the kernel of an operating system allows for multiple isolated user space instances, instead of just one. Such instances (often called containers, virtualization engines (VE), virtual private servers (VPS), or jails) may look and feel like a real server from the point of view of its owners and users.

On Unix-like operating systems, this technology can be seen as an advanced implementation of the standard chroot mechanism. In addition to isolation mechanisms, the kernel often provides resource management features to limit the impact of one container's activities on the other containers.

Contents [\[show\]](#)

Uses [\[edit\]](#)

Operating-system-level virtualization is commonly used in virtual hosting environments, where it is useful for securely allocating finite hardware resources amongst a large number of mutually-distrusting users. System administrators may also use it, to a lesser extent, for consolidating server hardware by moving services on separate hosts into containers on the one server.

Other typical scenarios include separating several applications to separate containers for improved security, hardware independence, and added resource management features. The improved security provided by the use of a chroot mechanism, however, is nowhere near ironclad.^[1] Operating-system-level virtualization implementations capable of live migration can also be used for dynamic load balancing of containers between nodes in a cluster.

Overhead [\[edit\]](#)

Operating-system-level virtualization usually imposes little to no overhead, because programs in virtual partitions use the operating system's normal system call interface and do not need to be subjected to emulation or be run in an intermediate virtual machine, as is the case with whole-system virtualizers (such as VMware ESXi, QEMU or Hyper-V) and paravirtualizers (such as Xen or UML). This form of virtualization also does not require support in hardware to perform efficiently.

Operating-system-level virtualization

From Wikipedia, the free encyclopedia

(Redirected from [Software container](#))

Operating-system-level virtualization is a server virtualization method where the kernel of an operating system allows for multiple isolated user space instances, instead of just one. Such instances (often called [containers](#), [virtualization engines](#) (VE), [virtual private servers](#) (VPS), or [jails](#)) may look and feel like a real server from the point of view of its owners and users.

On [Unix-like](#) operating systems, this technology can be seen as an advanced implementation of the standard [chroot](#) mechanism. In addition to isolation mechanisms, the kernel often provides resource management features to limit the impact of one container's activities on the other containers.

Contents [\[show\]](#)

Uses [\[edit\]](#)

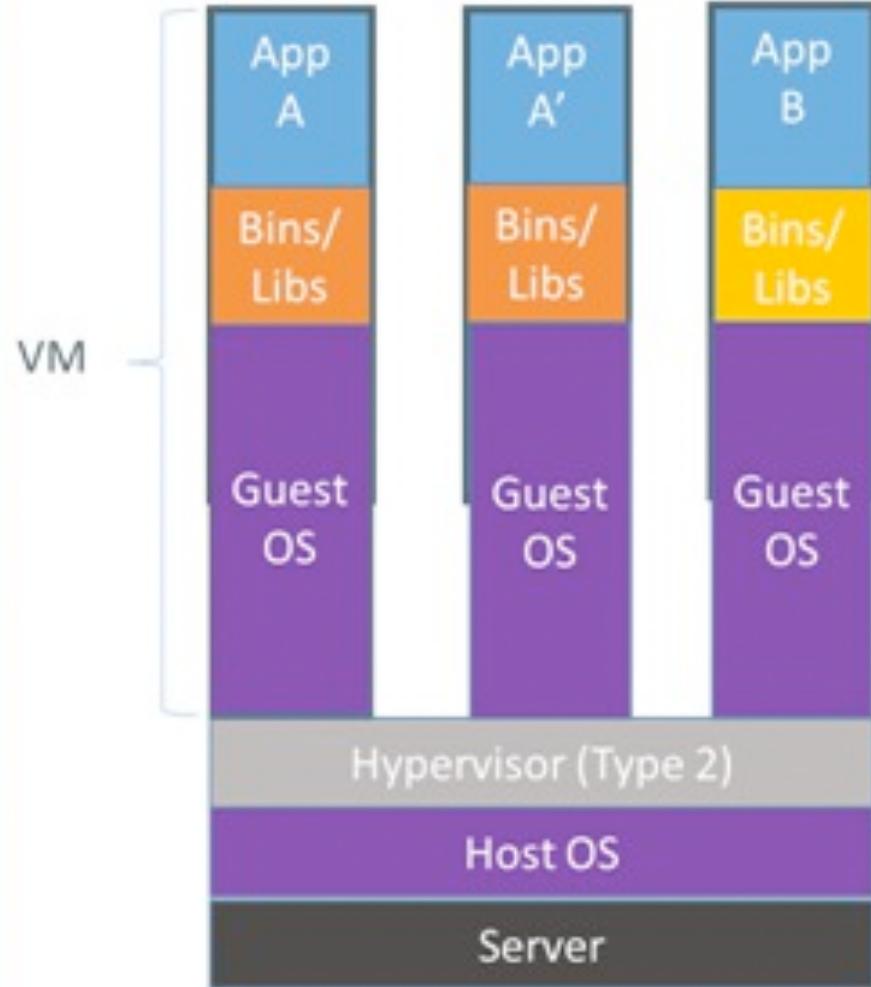
Operating-system-level virtualization is commonly used in [virtual hosting](#) environments, where it is useful for securely allocating finite hardware resources amongst a large number of mutually-distrusting users. System administrators may also use it, to a lesser extent, for consolidating server hardware by moving services on separate hosts into containers on the one server.

Other typical scenarios include separating several applications to separate containers for improved security, hardware independence, and added resource management features. The improved security provided by the use of a chroot mechanism, however, is nowhere near ironclad.^[1] Operating-system-level virtualization implementations capable of [live migration](#) can also be used for dynamic load balancing of containers between nodes in a cluster.

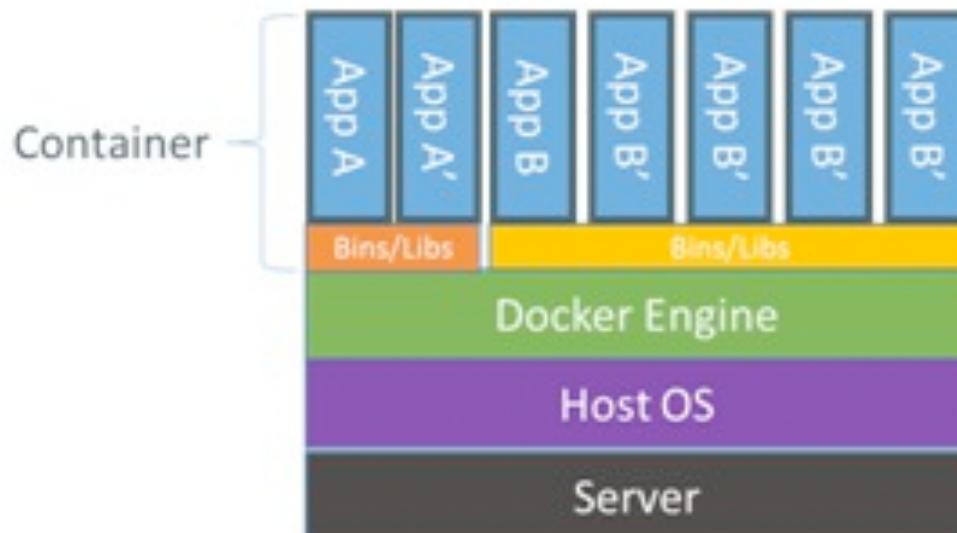
Overhead [\[edit\]](#)

Operating-system-level virtualization usually [imposes little to no overhead](#), because programs in virtual partitions use the operating system's normal [system call](#) interface and do not need to be subjected to [emulation](#) or be run in an intermediate [virtual machine](#), as is the case with whole-system virtualizers (such as [VMware ESXi](#), [QEMU](#) or [Hyper-V](#)) and paravirtualizers (such as [Xen](#) or [UML](#)). This form of virtualization also does not require support in hardware to perform efficiently.

Containers vs. VMs



Containers are isolated,
but share OS and, where
appropriate, bins/libraries



Layered approach

- Along other benefits, efficient:
- Docker pull brings deltas only

Writable Container

IMAGE

Add  mongoDB

IMAGE

Add  CouchDB

BASE IMAGE

Ubuntu

bootfs

cgroups, namespace, device mapper

Kernel

Docker ❤️ microservices

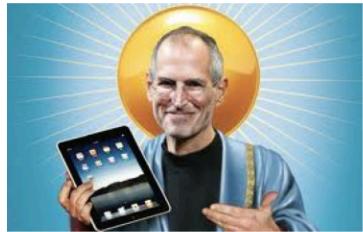
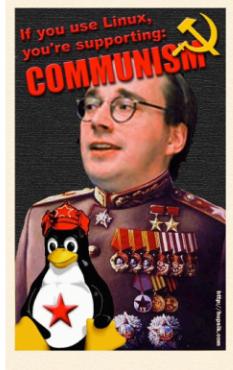
Docker fits much the micro-service arch of
bintray

(+ we do not need all services always)

Docker is not enough

Docker is not enough

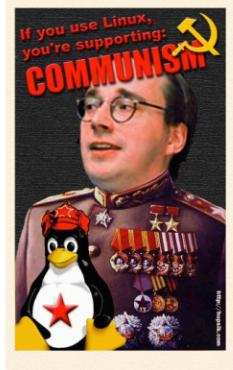
We are liberal



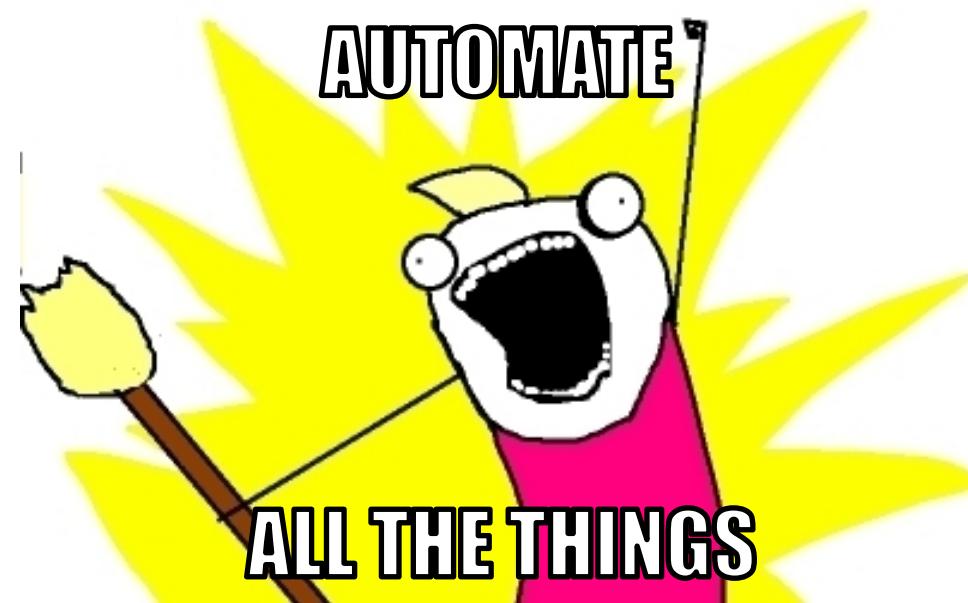
```
BANDLDR*** Address 0016950 has been at 8010000  
1.6.2 lsq!ifif SYSVER 0x00000000  
  
Name          Dll Base  StartEnd - Name  
ntoskrnl.exe  80010000  32157600  80100000  A.dll  
atapi.sys     80070000  33240000  80100000  S.I.P.R.T.  
disk.sys      801db000  33601500  80200000  A.B2.S.Y  
nfsa.sys      80237000  344mbw  80200000  A.B2.S.Y  
NTLice.sys    81148000  33ec6800  80100000  N.LP.P.Y  
turbo.sys     8228e000  31ec6690  80100000  T.U.B.O.S.Y  
rSecID.SYS    8230c000  338a0000  80100000  R.S.E.C.I.D.  
vihba.sys     82dca000  346mbw  80100000  V.I.H.B.A.  
zcdf.sys      82dc3000  346mbw  80100000  Z.C.D.F.  
schf.sys      82f16000  346mbw  80100000  S.C.H.F.  
sethc.sys     82f08000  346mbw  80100000  S.E.T.H.C.  
arpport.SYS   82f14000  346mbw  80100000  A.R.P.P.O.R.T.
```

Docker is not enough

We are liberal



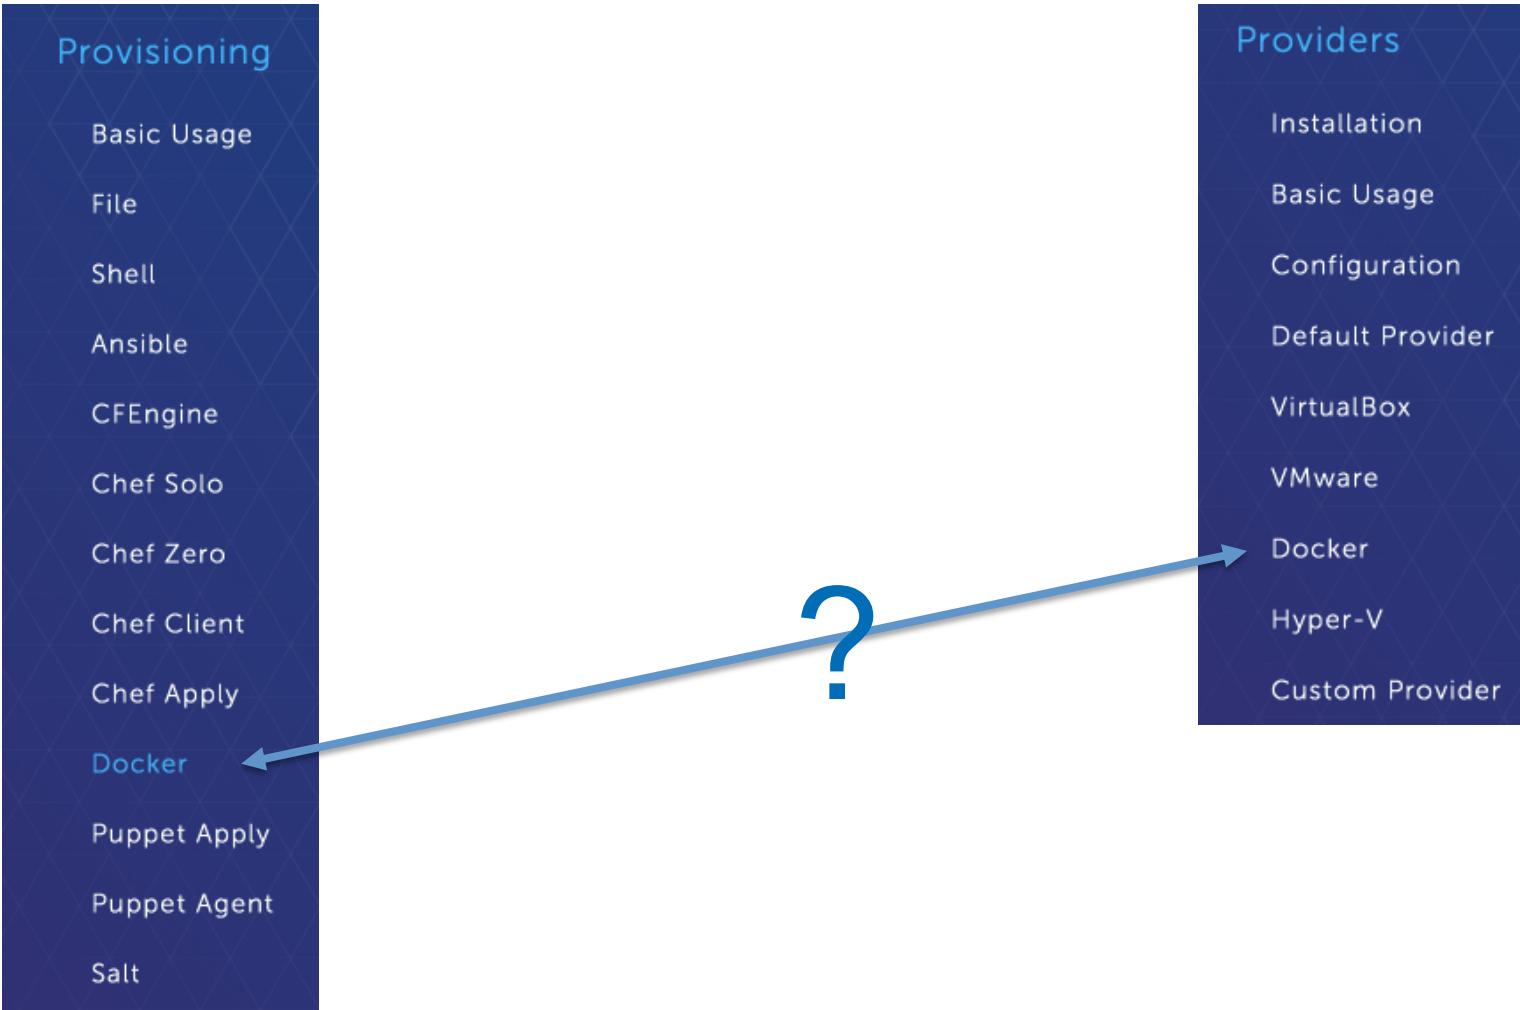
```
BANDLDR*** Address 0016950 has base at 8010000  
1.6.2 lsq!iflif SYSEVR 0x00000000  
  
Name          Dll Base  StartEnd - Name  
ntoskrnl.exe  80010000  32157600  80100000  A.dll  
atapi.sys     80070000  33240000  80100000  AIPRT.dll  
disk.sys      801db000  33601500  80100000  ARB2.SYS  
nfsa.sys      80237000  344mb0000  80100000  ARV1.SYS  
NTLice.sys    81148000  33ec6800  80100000  ARPV1.SYS  
turbo.sys     8228e000  31ec6690  80100000  ARV1.SYS  
rSecID.SYS   823c5000  33850000  80100000  ARV1.SYS  
vihba.sys     826c2000  3dc20000  80100000  ARV1.SYS  
zcdf.sys      82d35000  3dc35000  80100000  ARV1.SYS  
schf.sys      82f16000  3dc36000  80100000  ARV1.SYS  
sethc.sys    82f20000  3dc34000  80100000  ARV1.SYS  
parport.SYS  82f4d000  3dc34000  80100000  ARV1.SYS
```



Why Vagrant and Docker

- Configure host vm to run containers
 - Plugable Docker hosts
- Install Docker and run Docker
- Vagrant can control Docker containers –
 - run, destroy, link

Docker support in Vagrant



Provider? Provisioner?



Provider vs Provisioner

- Provider: defines the type of machines to manage, affects vm creation
- Provisioner: configures an existing machine on first vagrant up
- Vagrant destroy - undo machine creation

Solution II:

Vagrant + Docker
All the services in one container

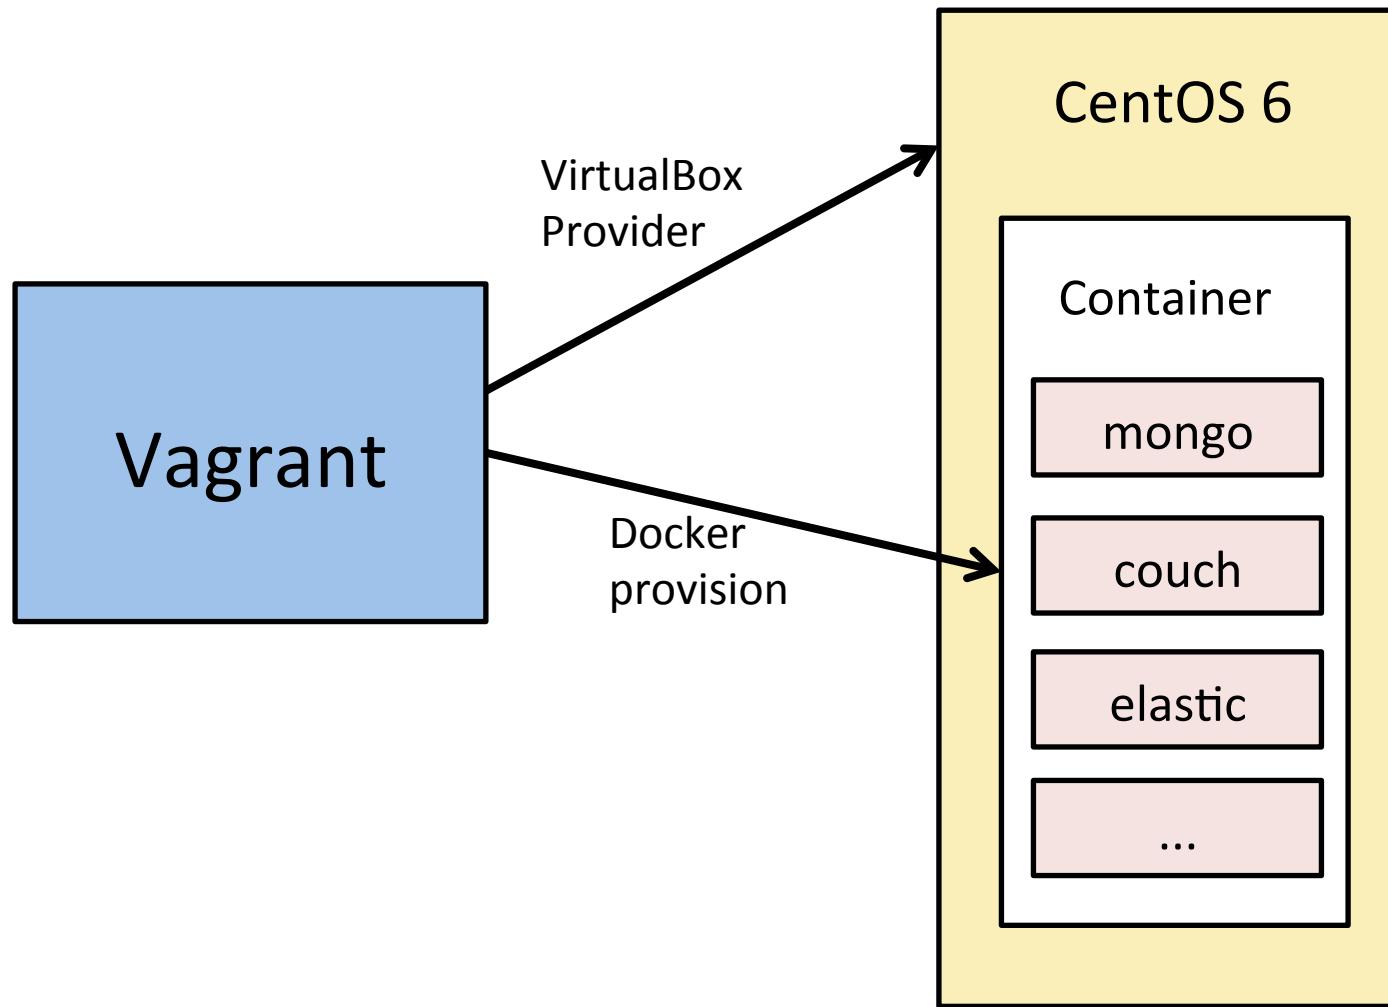
DOCKER PROVISIONER

Provisioner name: `"docker"`

The docker provisioner can automatically install [Docker](#), pull Docker containers, and configure certain containers to run on boot.

The docker provisioner is ideal for organizations that are using Docker as a means to distribute things like their application or services. Or, if you're just getting started with Docker

~~the Docker provisioner provides the easiest possible way to begin using Docker since the~~

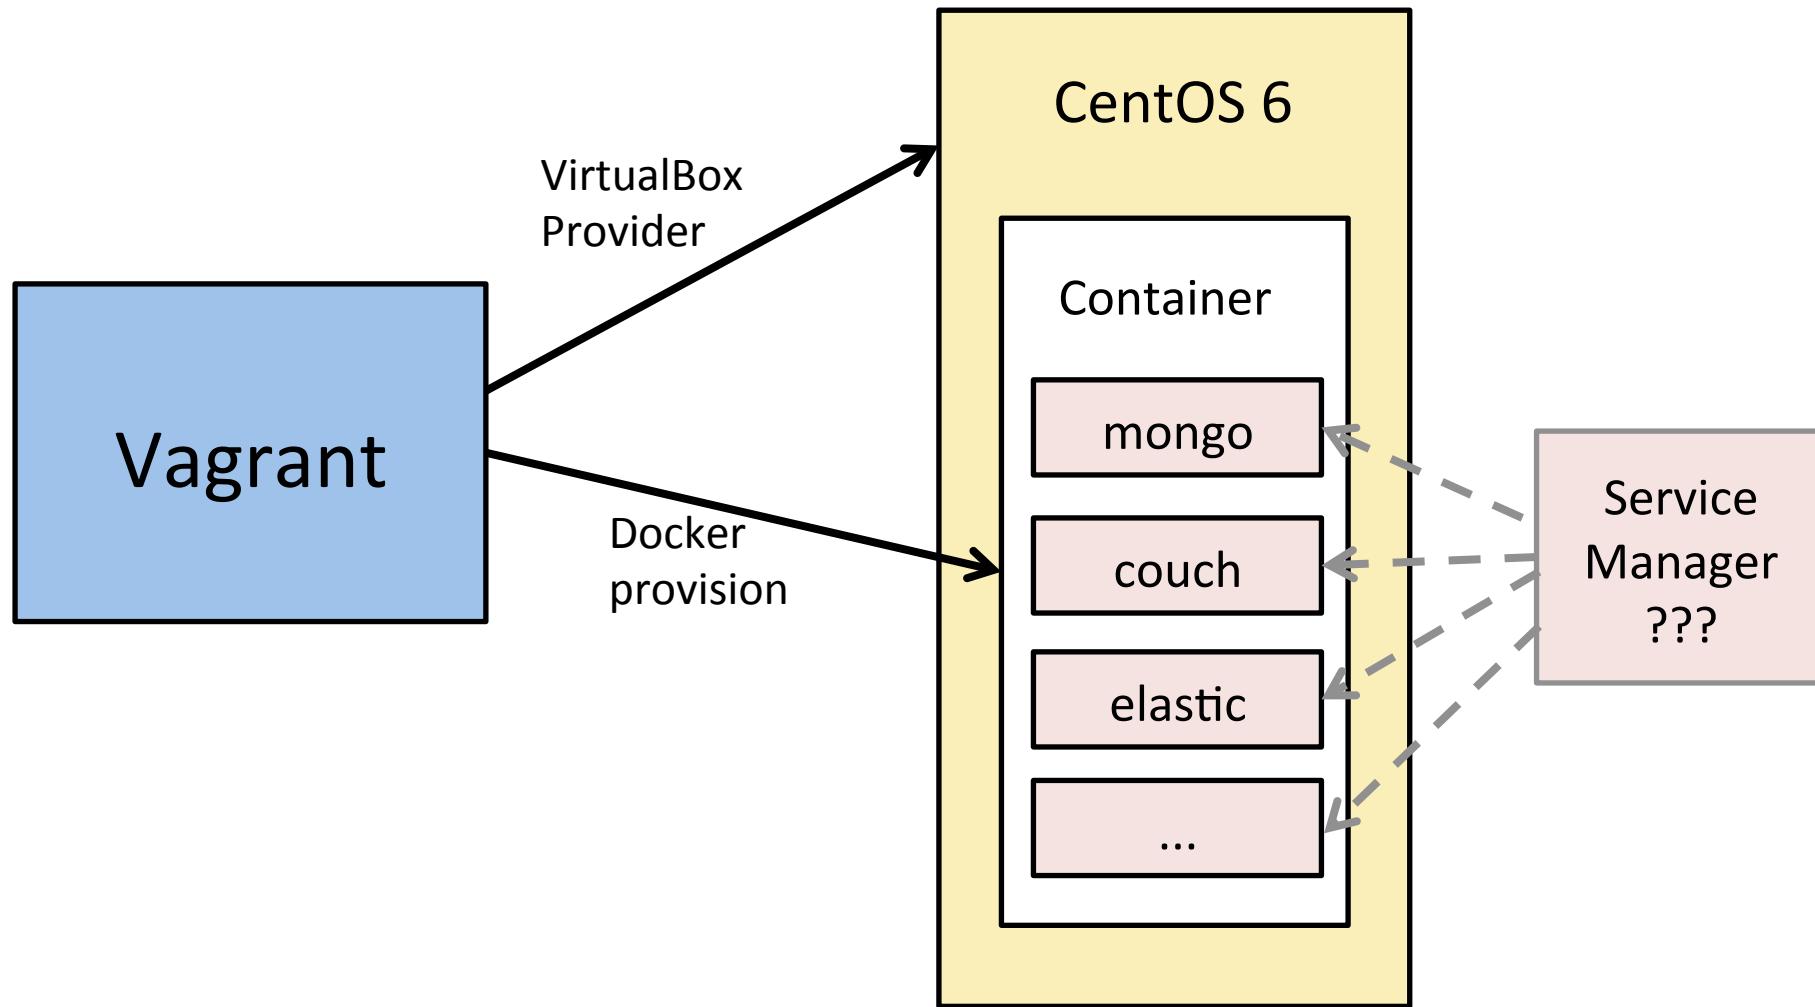


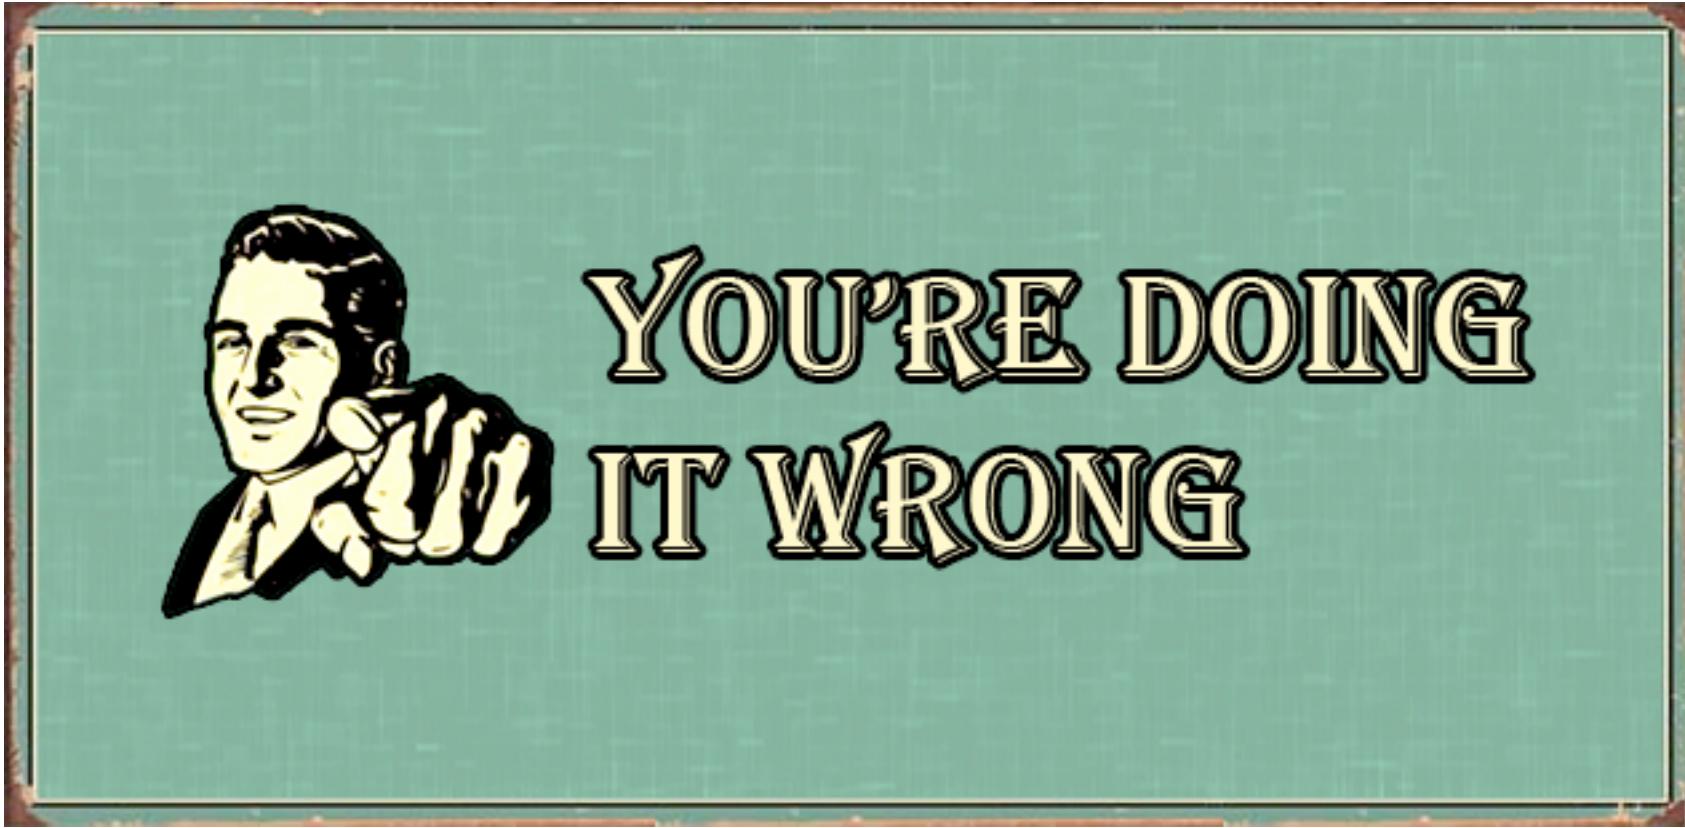
```
Vagrant.configure("2") do |config|
  config.vm.box = "centos6"

  config.vm.provider :virtualbox do |vb|
    vb.gui = false
  end

  ...

  config.vm.provision "docker" do |d|
    d.pull_images "jfrog/bintray"
    d.run "jfrog/bintray"
  end
end
```





- One single container or multiple containers - with one you also need to add a process manager; for instance Monit or Supervisor.

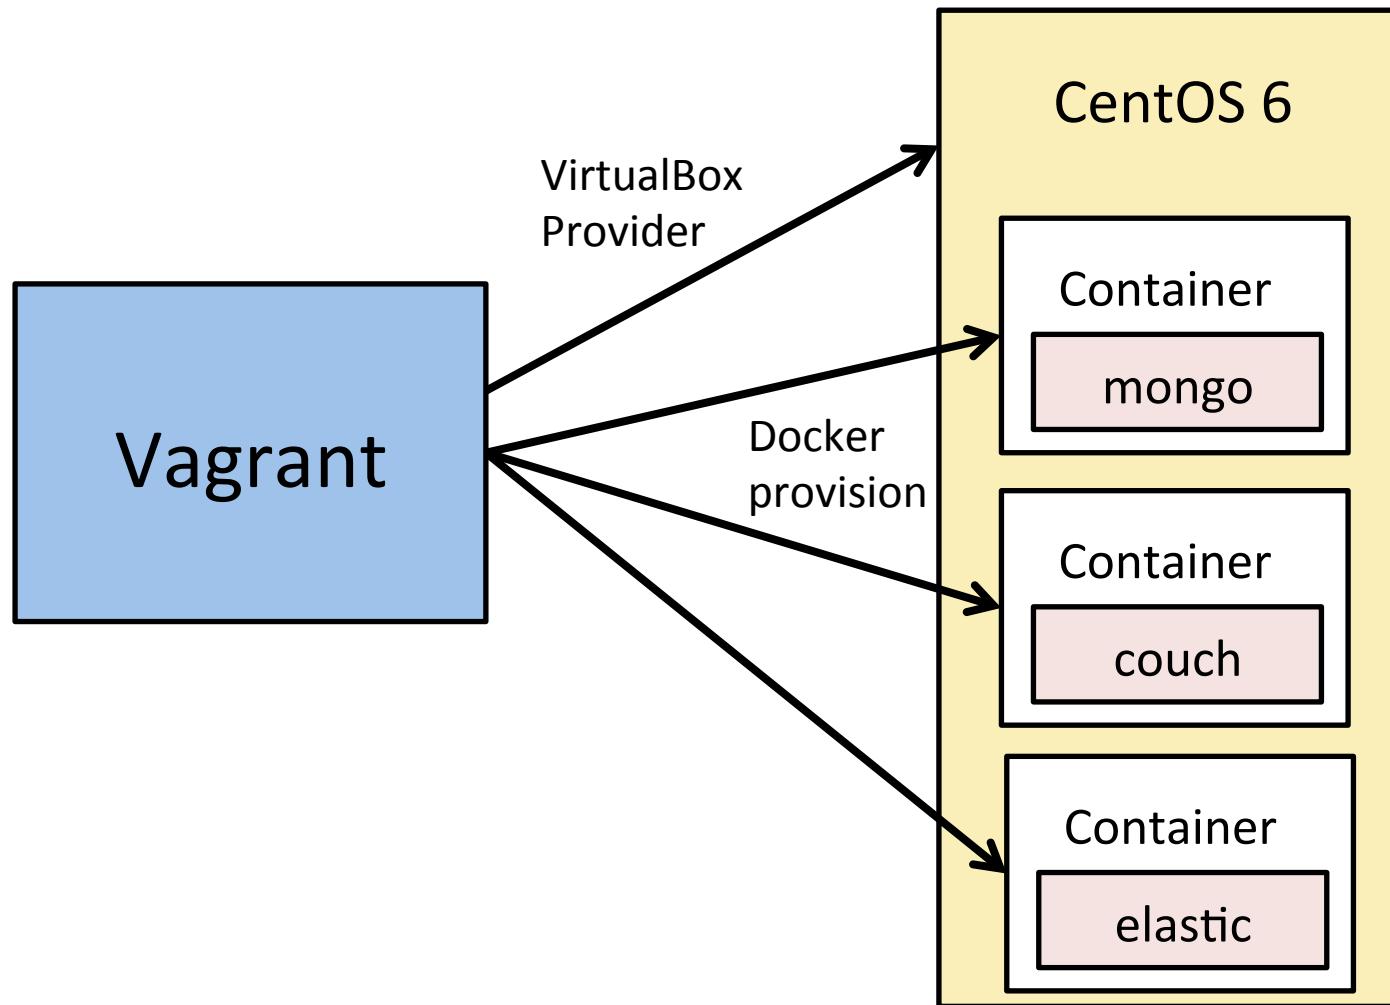
OK, we have an idea



Solution III:

Vagrant + Docker

Separate containers with Vagrant provisioners



```
Vagrant.configure("2") do |config|
  config.vm.box = "centos6"

  config.vm.provider :virtualbox do |vb|
    vb.gui = false
  end

  ...

  config.vm.provision "docker" do |d|
    d.pull_images "jfrog/mongo"
    d.run "jfrog/mongo"

    d.pull_images "jfrog/couch"
    d.run "jfrog/couch"
  end
end
```

But...



Developers break things



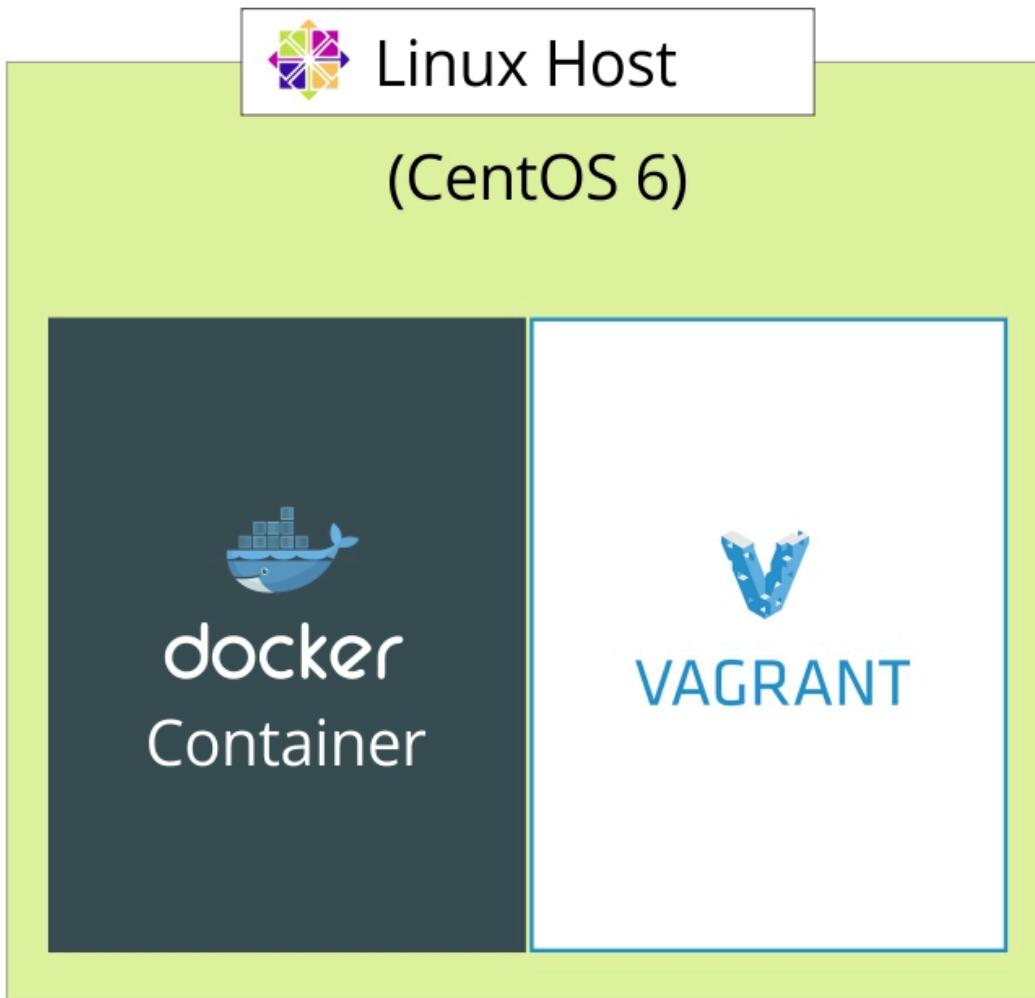
vagrant destroy

- Remember, vagrant destroy undoes the box completely
- I.E. empty box, no images, no configuration
- Next vagrant up will download all the images from scratch (specially painful over VPN)

Solution IV:

Vagrant + Docker
Separate containers with Vagrant providers
& use vagrant provider-management

Vagrant is smart!



Tiny Core Linux (a.k.a. boot2docker)

Boot2docker

Lightweight Linux for Docker

[Mac OS X](#)

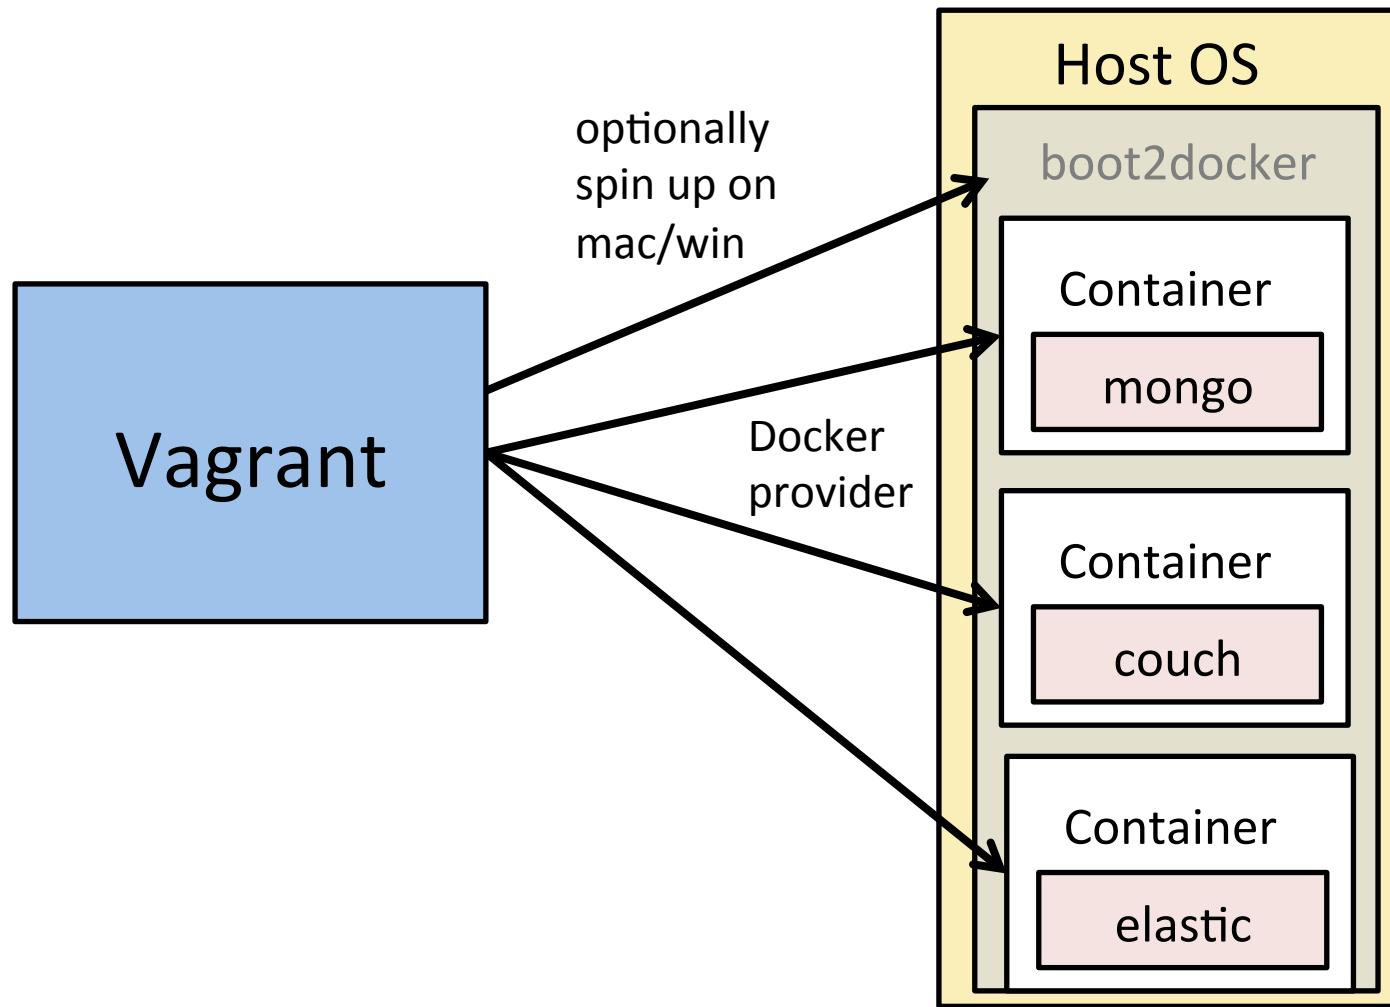
[Windows](#)

boot2docker

boot2docker is a lightweight Linux distribution based on [Tiny Core Linux](#) made specifically to run [Docker](#) containers. It runs completely from RAM, weighs ~27MB and boots in ~5s (YMMV).



The screenshot shows a Mac OS X software update window titled "Install Docker for Mac OS X". The window displays a progress bar and a message: "This package will run a program to determine if the software can be installed." Below the window, a portion of the Boot2Docker Mac OS X Install interface is visible, featuring a blue vertical bar and some text.



```
Vagrant.configure("2") do |config|
  //Look ma, no box!

  config.vm.define "mongodb" do |container|
    container.vm.provider "docker" do |d|
      ...
    end
  end

  config.vm.define "couch" do |container|
    container.vm.provider "docker" do |d|
      ...
    end
  end
end
```

How did we like it?

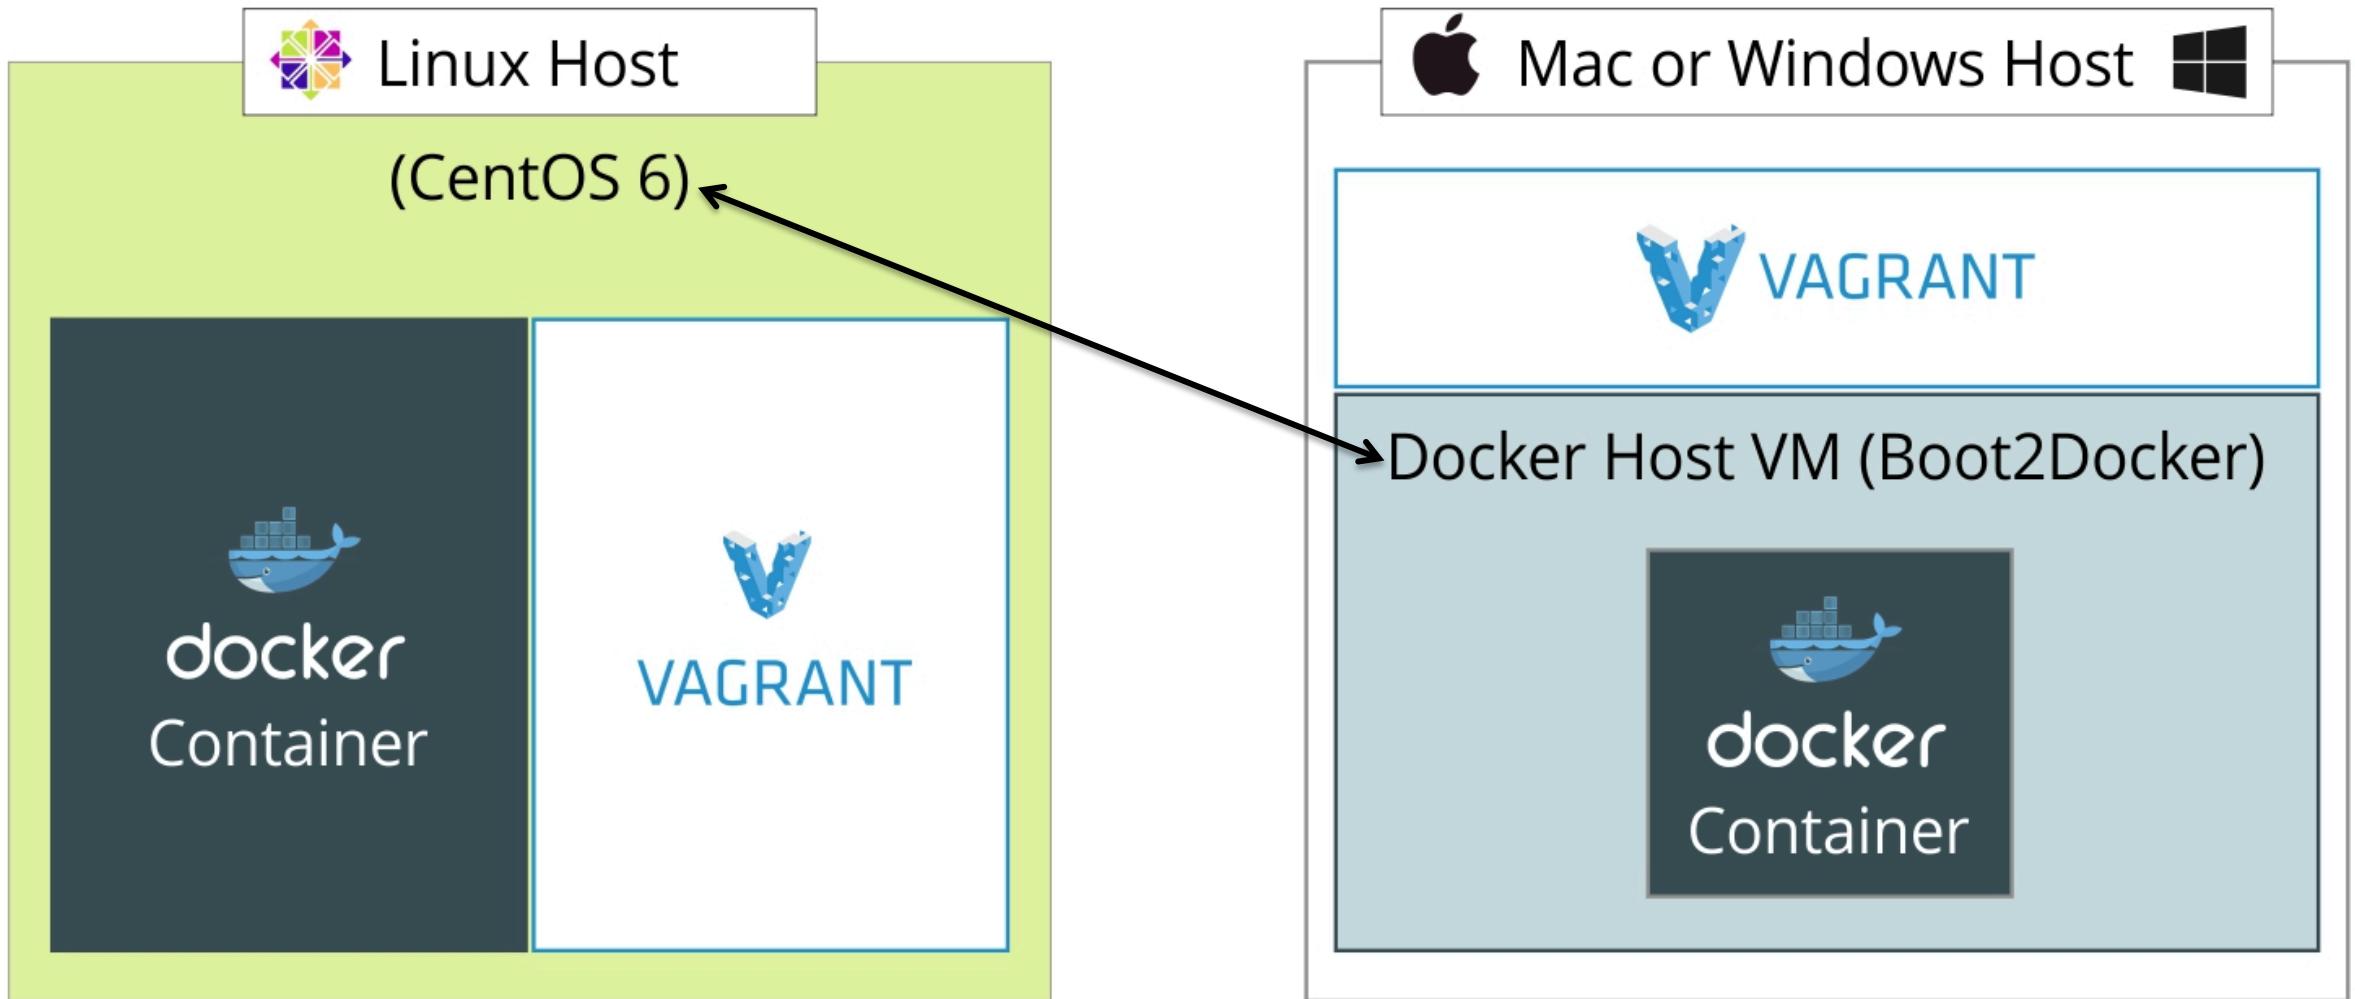


gifbox.ru

But...



The problem is right here:



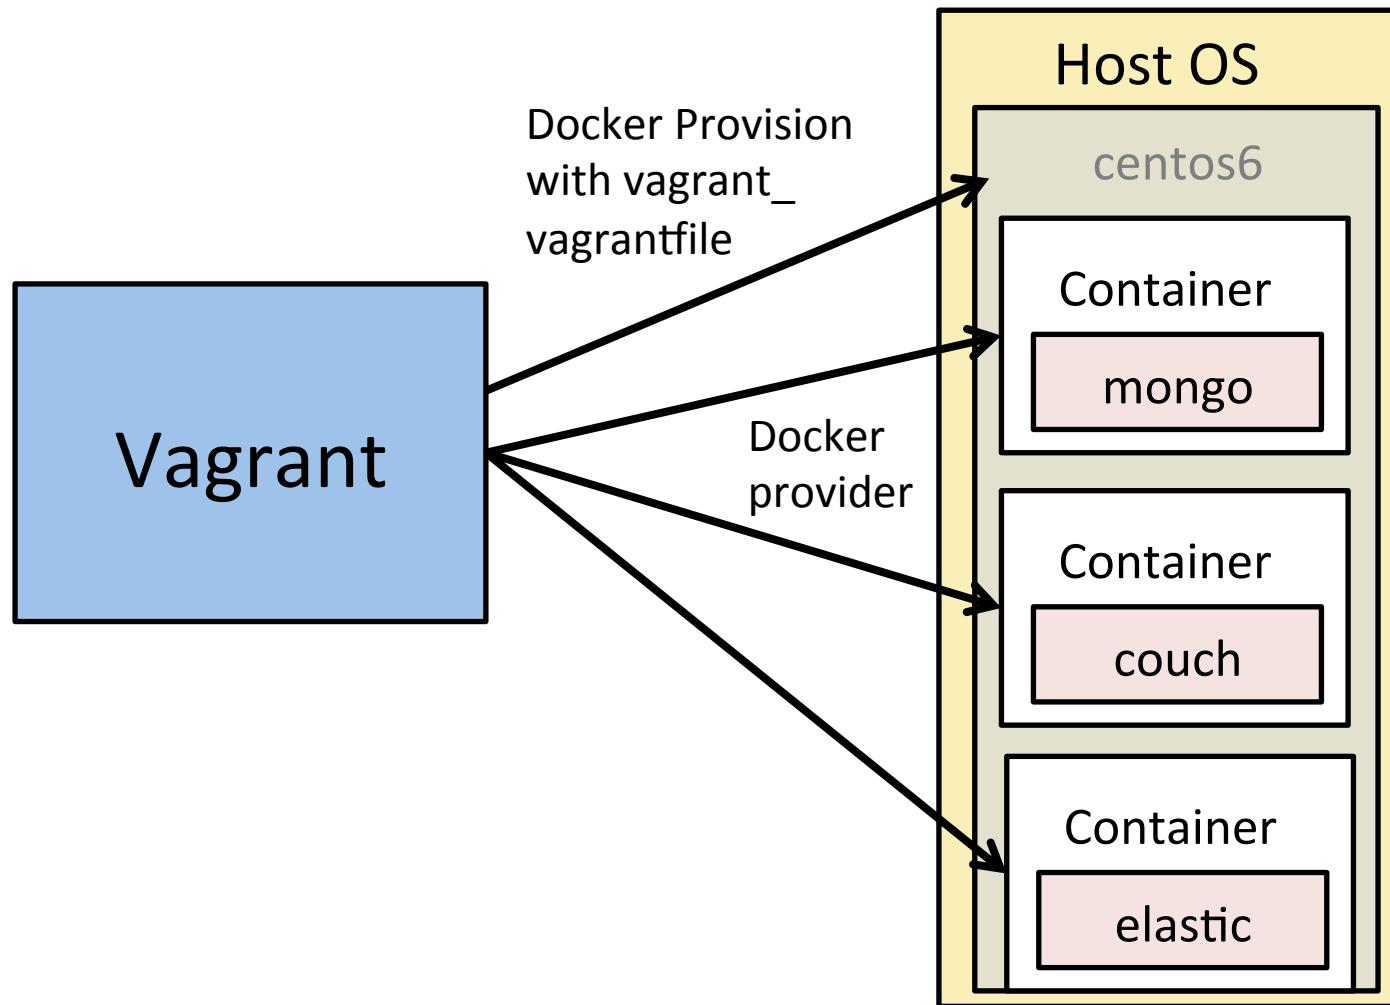
OK, we have an idea



Solution V:

Vagrant + Docker

Give up on boot2docker, use a box provider



```
Vagrant.configure("2") do |config|
  //Look ma, box is back!
  HOST_VAGRANTFILE = "#{checkout}/docker/Vagrantfile"

  config.vm.define "mongodb" do |container|
    container.vm.provider "docker" do |d|
      d.image = "#{repo_url}/bintray/mongo:2.6"
      d.vagrant_vagrantfile = "#{HOST_VAGRANTFILE}"
      ...
    end
  end

  config.vm.define "couch" do |container|
    container.vm.provider "docker" do |d|
      d.image = "#{repo_url}/bintray/couchdb:1.6.0"
      d.vagrant_vagrantfile = "#{HOST_VAGRANTFILE}"
      ...
    end
  end
end
```

```
Vagrant.configure("2") do |config|
  config.vm.provision "docker"

  config.vm.define "dockerhost"
  config.vm.box = "dockerhost"
  config.vm.box_url = "http://repo-local:8081/artifactory/bintray-
resolver/dockerhost-centos6.box"
  config.vm.box_download_checksum ="4ee5de076da2a4e1c7ecc7f201f408bb"
  config.vm.box_download_checksum_type = "md5"
  config.vm.synced_folder "../", "/vagrant", create: true
  config.vm.network "forwarded_port", guest: 5984, host: 5984
  ...
  config.vm.provider :virtualbox do |vb|
    vb.name = "dockerhost"
    vb.customize ["modifyvm", :id, "--cpus", 2, "--memory", 1512]
  end
end
```

How did we like it?



But...



Race condition on vagrant up

- When running multiple providers in the same vagrantfile, must use:
 --no-parallel
- SSHing to and destroying host vm is not trivial
- Devs started to get to know (and like) Docker

Official B2D is not up to date

- We want to use fork:

<https://github.com/YungSang/boot2docker-vagrant-box>

OK, we have an idea



Solution VI:

Vagrant + Docker

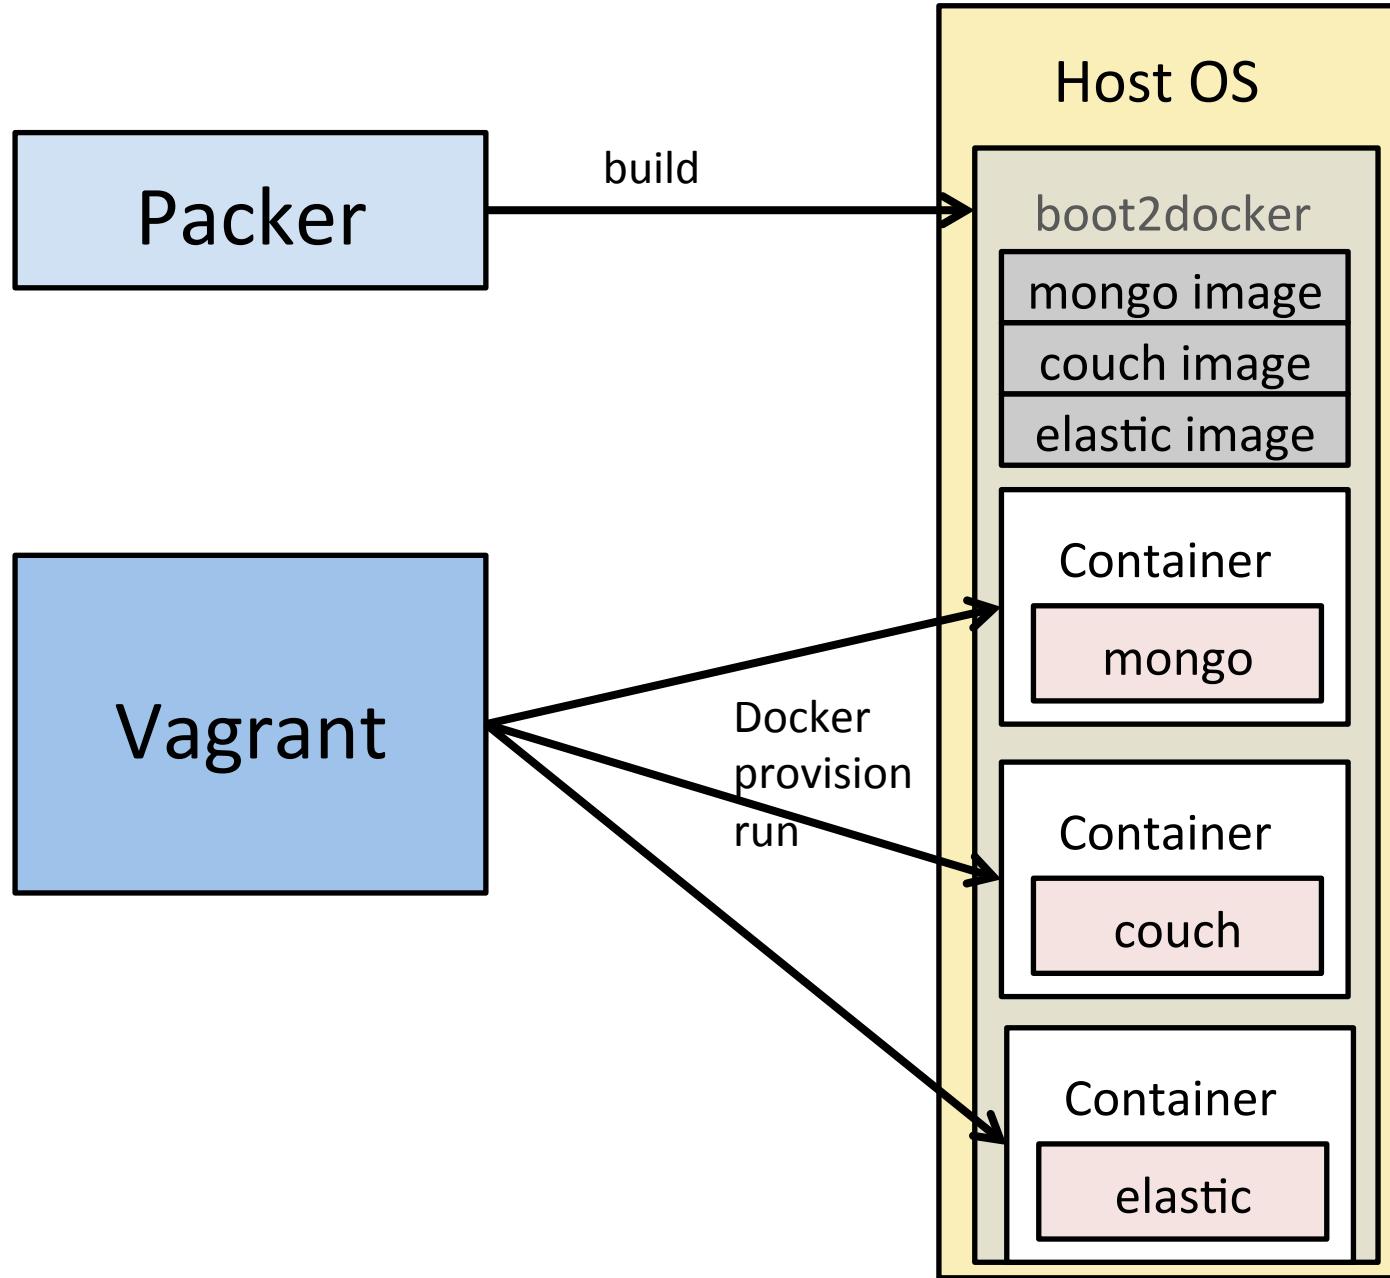
Full control: Make b2d and images part of the box

YungSang's boot2docker

- Well maintained – matches latest Docker versions
- Has ‘docker enter’ command – easy login to Docker

Build the VM with Packer

Packer is a tool
for creating
identical machine
images for
multiple
platforms from a
single source
configuration.



```
Vagrant.configure("2") do |config|
  config.vm.box = "bt-b2d-#{BOX_VERSION}"
  config.vm.box_url = "#{REPO}/bintray-boot2docker-vagrant/bt-
b2d-#{BOX_VERSION}.box"

  config.vm.provider :virtualbox do |vb, override|
    vb.gui = false
    ...
  end

  config.vm.provision :docker do |d|
    d.run "mongodb",
      image: "#{DOCKER_REPO}bintray/mongo:dev",
      args: "-p 27017:27017 -v #{HOST_DATA_FOLDER}/mongodb/
log/:/data/log/ --privileged=true",
      ...
    d.run "couch",
      image: "#{DOCKER_REPO}bintray/couchdb:dev",
      args: "-p 5984:5984 --sig-proxy=false -v
#{HOST_DATA_FOLDER}/couchdb/log/:/usr/local/var/log/couchdb/"
```

```
"provisioners": [
    ...
    {
        "type": "shell",
        "inline": [
            "docker pull repo:80/bintray/mongo:dev",
            "docker pull repo:80/bintray/redis:dev",
            "docker pull repo:80/bintray/couchdb:dev",
            "docker pull repo:80/bintray/elasticsearch:dev"
        ],
        "pause_before": "30s"
    }
],
```

How did we like it?



But...



Still, some challenges

- Updating the base box
 - Solved by timestamping
 - Deciding when to repackage the box
- Developers Docker learning curve

How it works

How it works

1. Packer builds custom b2d image with base images preinstalled

How it works

1. Packer builds custom b2d image with base images preinstalled
2. Vagrant downloads the box from artifactory

How it works

1. Packer builds custom b2d image with base images preinstalled
2. Vagrant downloads the box from artifactory
3. Vagrant runs the b2d virtual box

How it works

1. Packer builds custom b2d image with base images preinstalled
2. Vagrant downloads the box from artifactory
3. Vagrant runs the b2d virtual box
4. Vagrant pulls & starts the containers

How it works

1. Packer builds custom b2d image with base images preinstalled
2. Vagrant downloads the box from artifactory
3. Vagrant runs the b2d virtual box
4. Vagrant pulls & starts the containers
5. Vagrant destroy keeps Base containers

No, Thank you!

