

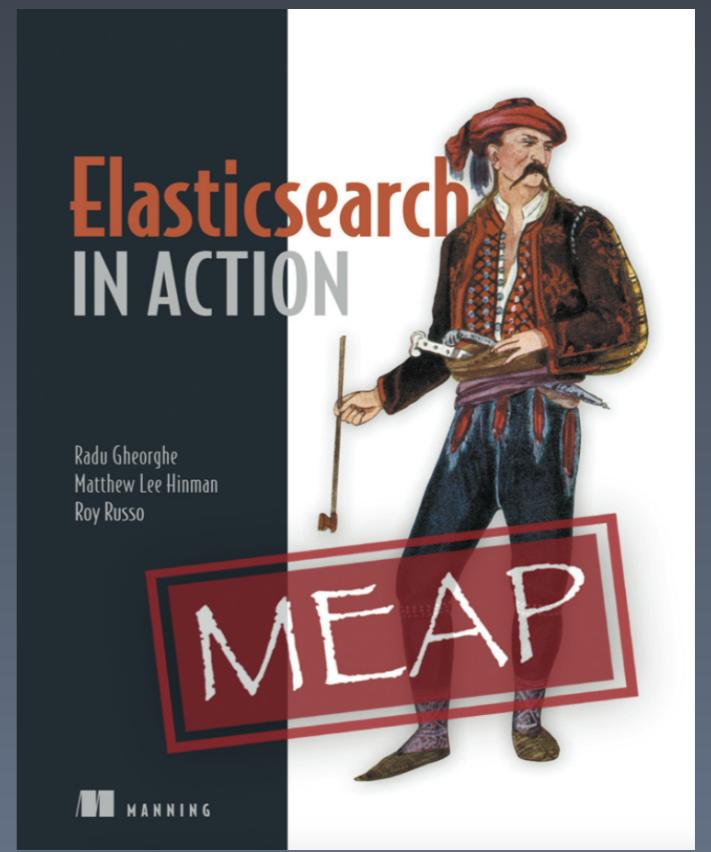
# Introduction To ElasticSearch

Real-Time Search and Analytics



# Who Am I

- Roy Russo
- VP Engineering, Predikto
- Co-Author - Elasticsearch in Action
  - Due ~April 2015.
- ElasticHQ.org
- Other (\*)



\* Silverpop, JBoss, AltiSource Labs

# Why Am I Here?

- What is Search
- What is Elasticsearch
- Real-World Use
- Scale Out
- Interacting with Elasticsearch

Search is about  
filtering  
information and  
determining  
relevance.

# How does a Search Engine

Select \* FROM make WHERE name LIKE '%Tesla  
%'

# Search Engines use Magic

## It's FM!

Where Magic == Inverted Index

# Inverted Index

- Take some documents
- Tokenize them
- Find unique tokens
- Map tokens to documents

	apple	oranges	peach
Document 1			
Document 2			
Document 3			
Document 4			
Document 5			
Document 6			

# Inverted Index

Search for “apple peach”

	apples	oranges	peach
Document 1			
Document 2			
Document 3			
Document 4			
Document 5			
Document 6			

# Relevance

- How many tokens per document?
- How many tokens relative to the number of total tokens in the document?
- What is the frequency of token across all documents?

# Relevance in Elasticsearch

- At Search Time
- At Index Time
- Term Frequency
  - Term / Document
- Inverse Document Frequency (IDF)
  - Term / All Documents in the collection
- Field-Length Norm

# What is Elasticsearch?

# Elasticsearch is...

- Search and Analytics engine
- Document Store
  - Every field is indexed/searchable
- Distributed

# What Elasticsearch is not

- Key-Value Store
  - Redis, Riak
- Column Family Store
  - C\*, HBase
- Graph Database
  - Neo4J



# ElasticSearch in a Nutshell

- Based on Apache Lucene
- Distributed
- Document-Oriented
- Schema free
- HTTP + JSON
- (Near) Real-time search
- Ecosystem
  - Hosting, Monitoring Apps, Clients (SDK)

# Where can I get it?

- Free and Open Source
- <https://www.elastic.co/>
- <https://github.com/elastic/elasticsearch>
- Backed by a Company, Elastic
  - Training
  - Support
  - Auth/AuthZ
  - Marvel for Monitoring

# How do I run it?

- Download it
  - <https://www.elastic.co/downloads>
- bin/elasticsearch
- <http://localhost:9200>

```
{  
  status: 200,  
  name: "Tesla",  
  cluster_name: "elasticsearch_royrusso",  
  version: {  
    number: "1.4.2",  
    build_hash: "927caff6f05403e936c20bf4529f144f0c89fd8c",  
    build_timestamp: "2014-12-16T14:11:12Z",  
    build_snapshot: false,  
    lucene_version: "4.10.2"  
  },  
  tagline: "You Know, for Search"  
}
```

# Elasticsearch requires Java. \*

\* You have 5 seconds to whine about it and then shutup.

# Some Use-Cases

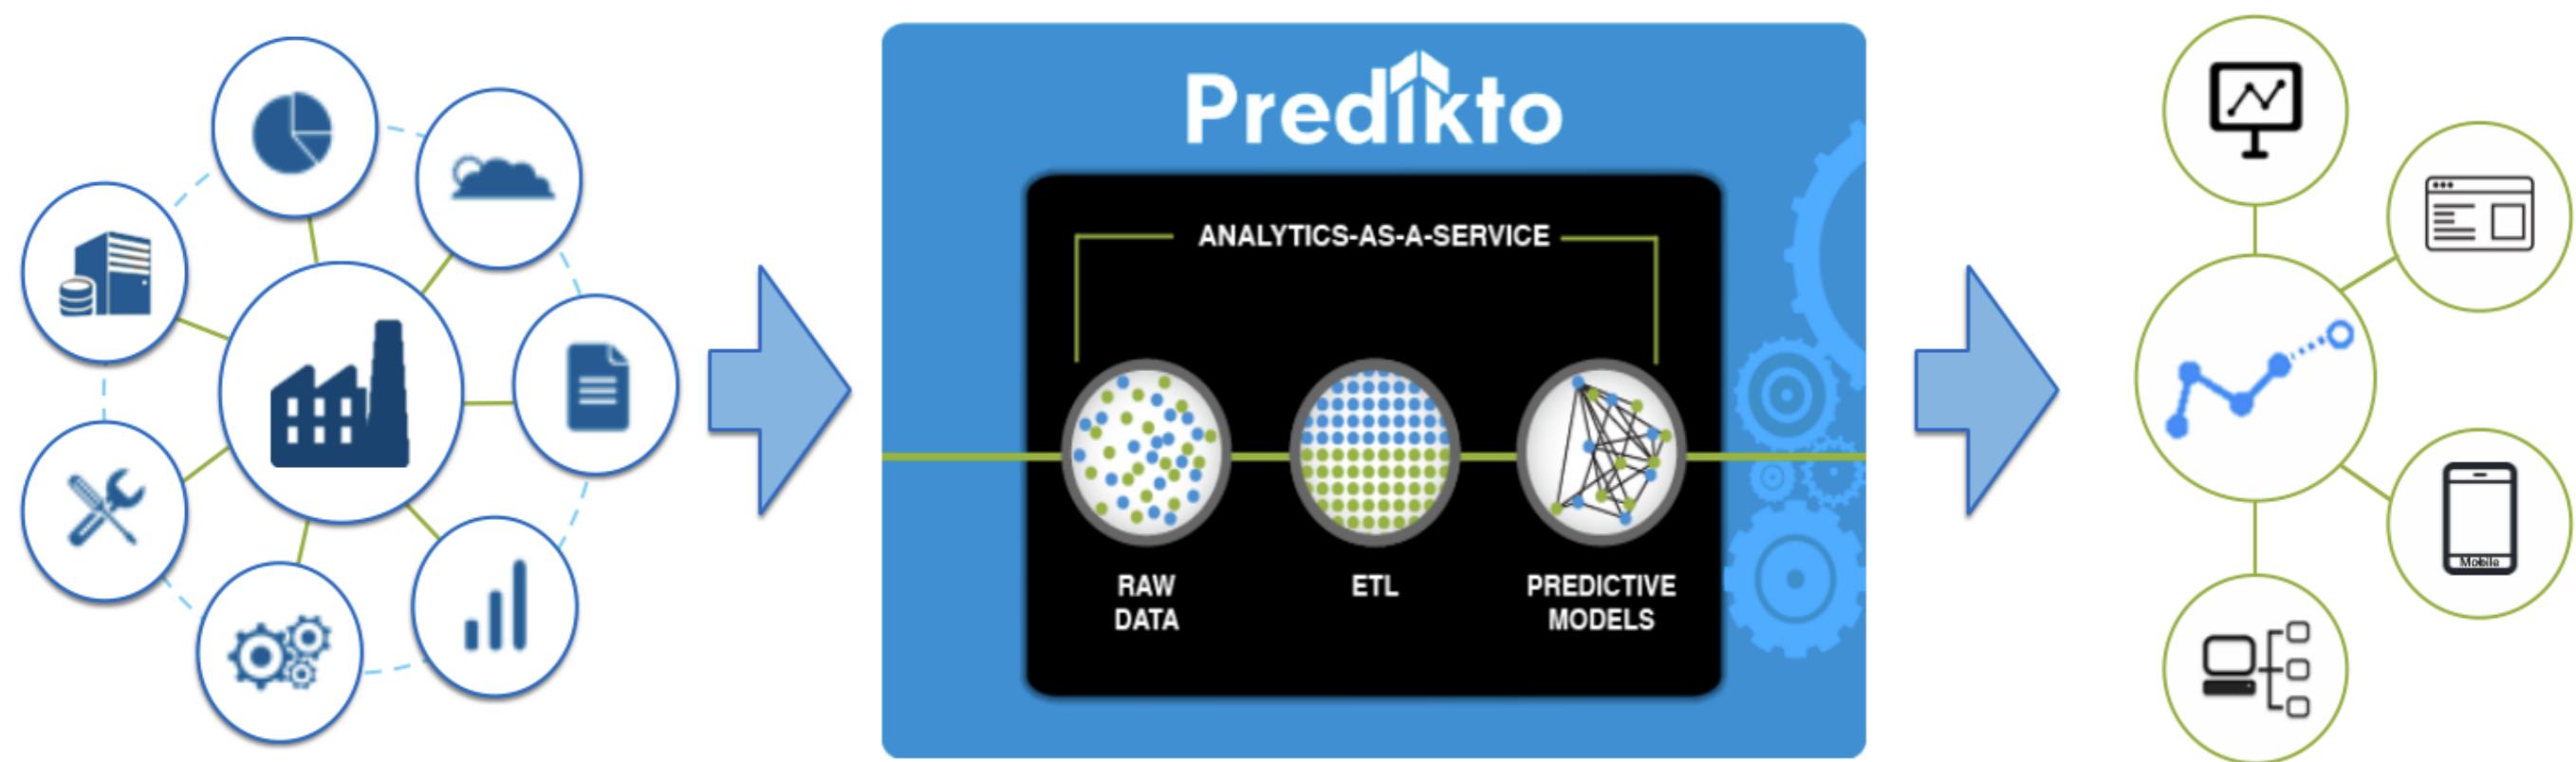
# ElasticSearch for Centralized

- Logstash + ElasticSearch + Kibana (ELK)
- Well... and then there's Luggly



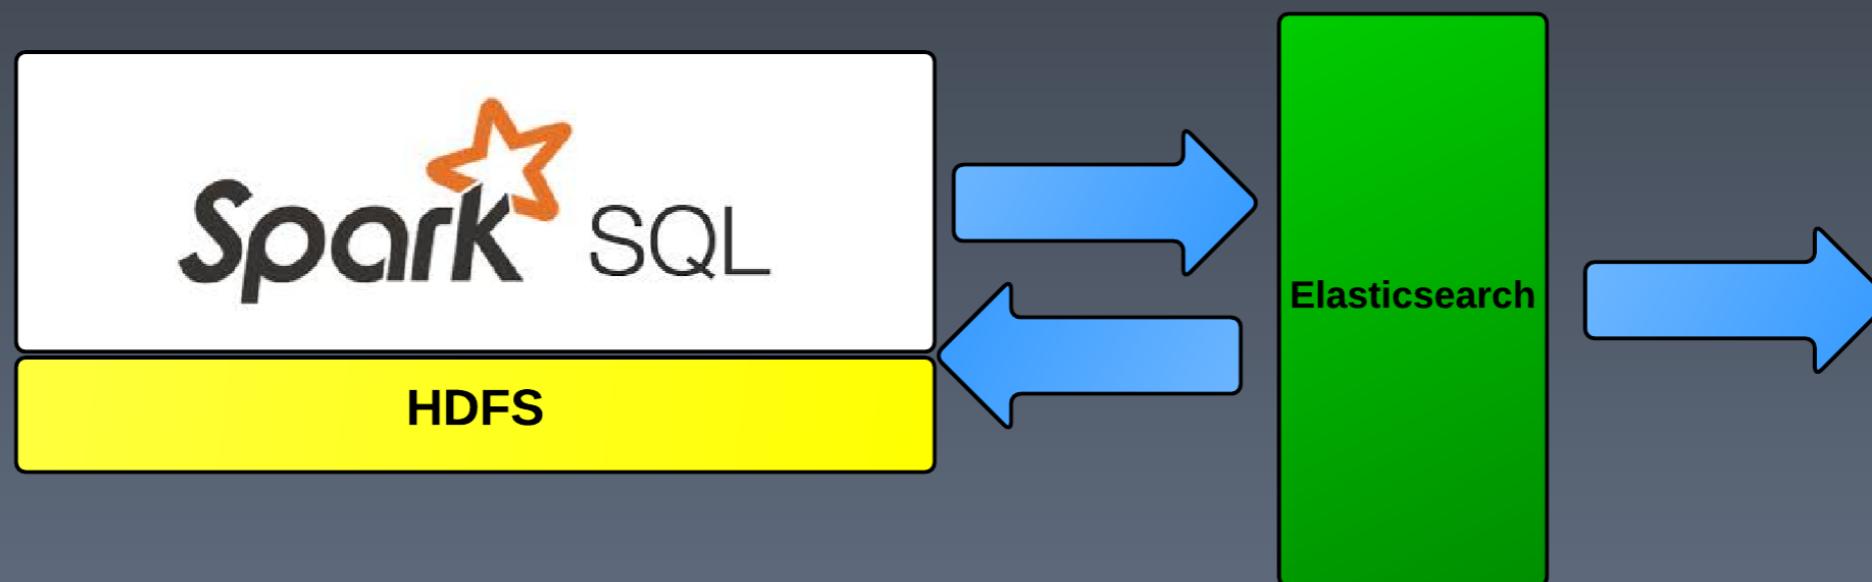
“Netflix is a Log generating company that happens to stream movies”

# Elasticsearch at Predikto



# Elasticsearch at Predikto

- Write From Spark to Elasticsearch
- Query from Spark to Elasticsearch
- Visualize



# Widely Used



NETFLIX



OpenTable™



# Based on Apache Lucene

- Free and Open Source
- Started in 1999
- Created by Doug Cutting
- What's it do?
  - Tokenizing
  - Locations
  - Relevance scoring
  - Filtering
  - Text search
  - Date Parsing



# Elasticsearch is a Document Store

# Document Store

- Like MongoDB and CouchDB
- Document DBs:
  - JSON documents
  - Collections of key-value collections
  - Nesting
  - Versioned

# What is a document?

```
{  
  "genre": "Crime",  
  "language": "English",  
  "country": "USA",  
  "runtime": 170,  
  "title": "Scarface",  
  "year": 1983  
}
```

# Modeled in JSON

```
{  
  "genre": "Crime",  
  "language": "English",  
  "country": "USA",  
  "runtime": 170,  
  "title": "Scarface",  
  "year": 1983  
}
```



```
{  
  "_index": "imdb",  
  "_type": "movie",  
  "_id": "u17o8zy9RcKg6SjQZqQ40w",  
  "_version": 1,  
  "_source": {  
    "genre": "Crime",  
    "language": "English",  
    "country": "USA",  
    "runtime": 170,  
    "title": "Scarface",  
    "year": 1983  
  }  
}
```

# Schema-Free

- Dynamic Mapping
  - Elasticsearch guesses the data-types (string, int, float...)

```
"imdb": {  
    "movie": {  
        "properties": {  
            "country": {  
                "type": "string",  
                "store":true,  
                "index":false  
            },  
            "genre": {  
                "type": "string",  
                "null_value" : "na",  
                "store":false,  
                "index":true  
            },  
            "year": {  
                "type": "long"  
            }  
        }  
    }  
}
```

# Elasticsearch is Distributed

# Terminology

MySQL	Elasticsearch
Database	Index
Table	Type
Row	Document
Column	Field
Schema	Mapping
Index	(Everything is indexed)
SQL	Query DSL

- Cluster: 1..N Nodes w/ same Cluster Name
- Node: One ElasticSearch instance (1 java proc)
- Shard = One Lucene instance
  - 0 or more replicas

# High Availability

- No need for load balancer
- Different Node Types
- Indices are Sharded
- Replica shards on different Nodes
- Automatic Master election & failover

# About Indices / Shards

```
$ curl -XPUT 'http://localhost:9200/twitter/' -d '{  
  "settings" : {  
    "index" : {  
      "number_of_shards" : 3,  
      "number_of_replicas" : 2  
    }  
  }'  
}'
```

# Cluster Topology

4 Node Cluster



Index A: 2 Shards & 1 Replica

Index B: 3 Shards & 1 Replica

# Discovery

- Nodes discover each other using multicast.
  - Unicast is an option
- Each cluster has an elected master node
  - Beware of split-brain

```
discovery.zen.ping.multicast.enabled: false
discovery.zen.ping.unicast.hosts: ["host1", "host2:port", "host3"]
```

# Nodes

- Master node handles cluster-wide (Meta-API) events:
  - Node participation
  - New indices create/delete
  - Re-Allocation of shards
- Data Nodes
  - Indexing / Searching operations
- Client Nodes
  - REST calls
  - Light-weight load balancers

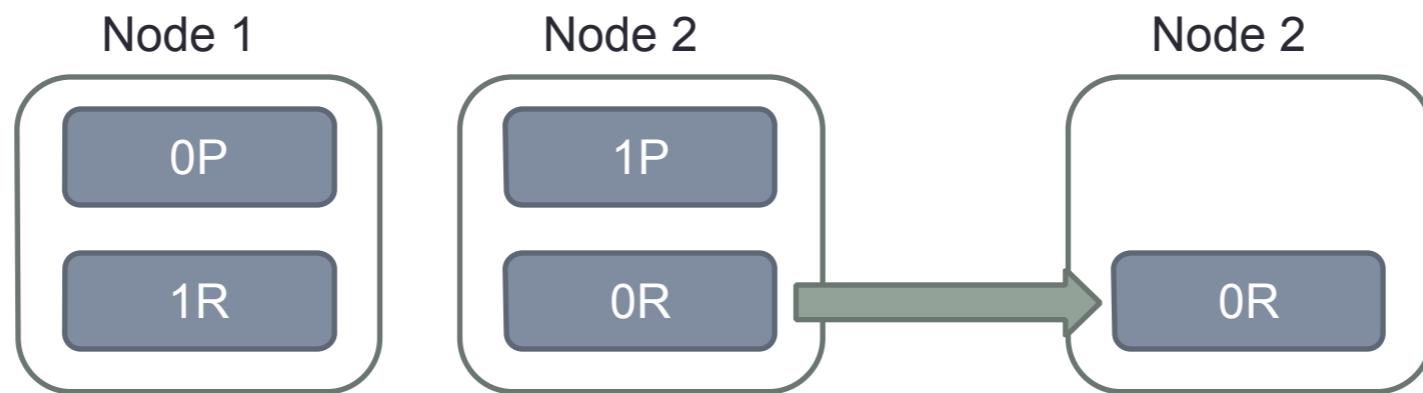
node.data | node.master

# The Basics - Shards

- Primary Shard:
  - First time Indexing
  - Index has 1..N primary shards (default: 5)
  - # Not changeable once index created
- Replica Shard:
  - Copy of the primary shard
  - # Can be changed later
  - Each primary has 0..N replicas
  - HA:
    - Promoted to primary if primary fails

# Shard Auto-Allocation

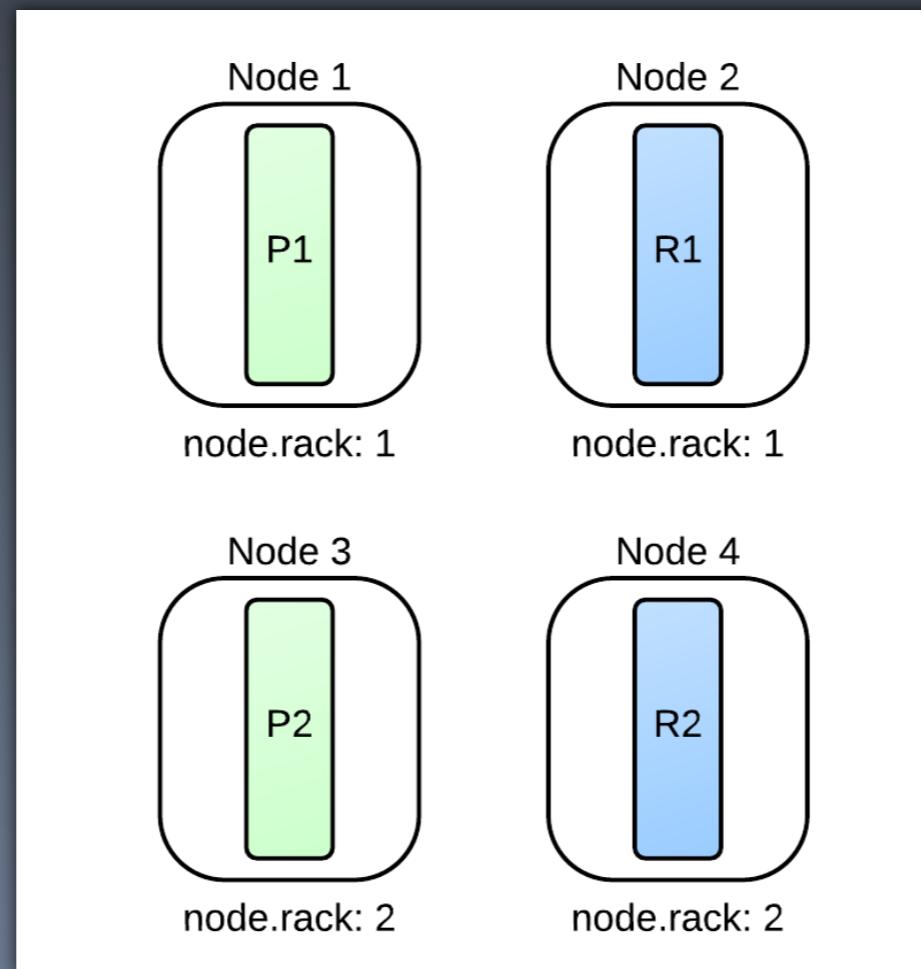
## Add a Node: Shards Relocate



- Shard Phases:
  - Unassigned
  - Initializing
  - Started
  - Relocating

# Allocation Awareness

- Shard Allocation Awareness
  - `cluster.routing.allocation.awareness.attributes: rack`
  - Shards RELOCATE to even distribution
  - Primary & Replica will NOT be on the same rack



# Cluster State

- Cluster State

- Node Membership
- Indices Settings
- Shard Allocation Table

```
cURL -XGET http://localhost:9200/_cluster/state?pretty=1
{
  "cluster_name" : "elasticsearch_royrusso",
  "version" : 27,
  "master_node" : "s3fpXfPKSFeUqo1MYZxSng",
  "blocks" : { },
  "nodes" : {
    "s3fpXfPKSFeUqo1MYZxSng" : {
      "name" : "Bulldozer",
      "transport_address" : "inet[localhost/127.0.0.1:9300]",
      "attributes" : { }
    }
  },
  "metadata" : {
    "templates" : {
      "logging_index_all" : {
        "template" : "logstash-09-*",
        "order" : 1,
        "settings" : {
          "index" : {
            "number_of_shards" : "2",
            "number_of_replicas" : "1"
          }
        },
        "mappings" : {
          "date" : {
            "store" : false
          }
        }
      },
      "logging_index" : {
        "template" : "logstash-*",
        "order" : 0,
        "settings" : {
          "index" : {
            "number_of_shards" : "2",
            "number_of_replicas" : "1"
          }
        },
        "mappings" : {
          "date" : {
            "store" : true
          }
        }
      }
    }
  }
}
```

# Talking to Elasticsearch

# REST

- HTTP Verbs: GET, POST, PUT, DELETE
- JSON
- \_cat API

```
% curl '192.168.56.10:9200/_cat/health?v&ts=0'  
cluster status nodeTotal nodeData shards pri relo init unassign  
foo      green          3        3     3   3    0    0      0
```

# The API

- Document
- Cluster
  - Node
- Index
- Search

# Create a Document

```
curl -XPOST 'http://127.0.0.1:9200/imdb/movie' -d '  
{  
    "genre": "Crime",  
    "language": "English",  
    "country": "USA",  
    "runtime": 170,  
    "title": "Scarface",  
    "year": 1983  
}';
```

```
{  
    "_index":"imdb",  
    "_type":"movie",  
    "_id":"AUwGeWib1u4mCngDYT7y",  
    "_version":1,  
    "created":true  
}
```

# Of note...

```
curl -XPOST 'http://127.0.0.1:9200/imdb/movie' -d '  
{  
    "genre": "Crime",  
    "language": "English",  
    "country": "USA",  
    "runtime": 170,  
    "title": "Scarface",  
    "year": 1983  
}';
```

Auto-creates Index & Type

Auto-Gen ID

```
{  
    "_index":"imdb",  
    "_type":"movie",  
    "_id":"AUwGeWib1u4mCngDYT7y",  
    "_version":1,  
    "created":true  
}
```

Auto-Version

# Get a Document

```
curl -XGET 'http://127.0.0.1:9200/imdb/movie/AUwGeWib1u4mCngDYT7y';
```

```
{
  "_index": "imdb",
  "_type": "movie",
  "_id": "AUwGeWib1u4mCngDYT7y",
  "_version": 1,
  "found": true,
  "_source":
  {
    "genre": "Crime",
    "language": "English",
    "country": "USA",
    "runtime": 170,
    "title": "Scarface",
    "year": 1983
  }
}
```

# Update a Document

```
curl -XPUT 'http://127.0.0.1:9200/imdb/movie/AUwGeWib1u4mCngDYT7y' -d '  
{  
    "genre": "Crime",  
    "language": "English",  
    "country": "USA",  
    "runtime": 180,  
    "title": "Scarface",  
    "year": 1983  
}';
```

```
{  
    "_index":"imdb",  
    "_type":"movie",  
    "_id":"AUwGeWib1u4mCngDYT7y",  
    "_version":2,  
    "created":false  
}
```

\* More like an Upsert.

# Delete a Document

```
curl -XDELETE 'http://127.0.0.1:9200/imdb/movie/AUwGeWib1u4mCngDYT7y'
```

```
{  
  "found":true,  
  "_index":"imdb",  
  "_type":"movie",  
  "_id":"AUwGeWib1u4mCngDYT7y",  
  "_version":2  
}
```

# You can also...

- Partial document updating
- Specify Version
- Specify ID
- Multi-Get API
- Exists API
- Bulk API

# How Searching Works

- How it works:
  - Search request hits a node
  - Node broadcasts to every shard in the index
  - Each shard performs query
  - Each shard returns metadata about results
  - Node merges results and scores them
  - Node requests documents from shards
  - Results merged, sorted, and returned to client.

# REST API - Search

- Free Text Search
  - URL Request
- Complex Query

```
http://localhost:9200/imdb/movie/_search?q=scar*
```

```
http://localhost:9200/imdb/movie/_search?q=scarface+OR+star
```

```
http://localhost:9200/imdb/movie/_search?q=(scarface+OR+star)+AND+year:[1981+TO+1984]
```

# REST API – Query DSL

```
curl -XPOST 'localhost:9200/_search?pretty' -d '{
  "query" : {
    "bool" : {
      "must" : [
        {
          "query_string" : {
            "query" : "scarface or star"
          }
        },
        {
          "range" : {
            "year" : { "gte" : 1931 }
          }
        }
      ]
    }
  }
}'
```

# REST API – Query DSL

- Boolean Query

```
"bool":{  
    "must": [  
        {  
            "match":{  
                "color":"blue"  
            }  
        },  
        {  
            "match":{  
                "title":"shirt"  
            }  
        }  
    ],  
    "must_not": [  
        {  
            "match":{  
                "size":"xxl"  
            }  
        }  
    ],  
    "should": [  
        {  
            "match":{  
                "textile":"cotton"  
            }  
        }  
    ]  
}
```

# REST API – Query DSL

- Range Query
  - Numeric / Date Types
- Prefix/Wildcard Query
  - Match on partial terms
- RegExp Query
- Geo\_bbox
  - Bounding box filter
- Geo\_distance
  - Geo\_distance\_range

```
{  
  "range":{  
    "founded_year":{  
      "gte":1990,  
      "lt":2000  
    }  
  }  
}
```

# Filters

- Filters recommended over Queries
  - Better cache support

```
curl -XGET 'http://localhost:9200/my_index/events/_search?pretty=1' -d '  
{  
  "from" : 0,  
  "size" : 0,  
  "query" : {  
    "terms" : {  
      "message" : [ "apples" ],  
      "minimum_should_match" : "1"  
    }  
  },  
  "post_filter" : {  
    "terms" : {  
      "userId" : [ "25476c6788ce", "g20d5470d7b4" ],  
      "execution" : "or"  
    }  
  },  
  "sort": { "eventDate": { "order": "desc" }},  
  "explain" : false  
};
```

# Analyzers / Tokenizers

```
curl -XPUT 'http://localhost:9200/my_index/' -d '  
{  
  "settings" :  
  {  
    "analysis" : {  
      "analyzer" : {  
        "str_search_analyzer" : {  
          "tokenizer" : "keyword",  
          "filter": ["lowercase"]  
        },  
        "str_index_analyzer" : {  
          "tokenizer" : "substring",  
          "filter" : ["lowercase", "stop"]  
        }  
      },  
      "tokenizer" : {  
        "substring" : {  
          "type" : "edgeNgram",  
          "min_gram" : "3",  
          "max_gram" : "42",  
          "token_chars" : ["letter", "digit"]  
        }  
      }  
    }  
}
```

```
curl -XPUT 'http://localhost:9200/my_index/events/_mapping' -d '  
{  
  "events" : {  
    "properties" :  
    {  
      "eventId" : {"type" : "string", "store" : true, "index" : "not_analyzed" },  
      "userId" : {"type" : "string", "store" : false, "index" : "not_analyzed" },  
      "message" : {  
        "type" : "string", "store" : false,  
        "search_analyzer" : "str_search_analyzer",  
        "index_analyzer" : "str_index_analyzer"  
      }  
    }  
  }  
};
```

# Tokenizers

- Whitespace
- NGram
- Edge NGram
- Letter
  - @ non-letters

# Clients

- Client list: <http://www.elasticsearch.org/guide/clients/>
  - Java (Node) Client, JS, PHP, Perl, Python, Ruby
- Spring Data:
  - Uses TransportClient
  - Implementation of ElasticsearchRepository aligns with generic Repository interfaces

# Monitoring

- BigDesk
- Kopf
- Head
- ElasticHQ
- Marvel
- Sematext SPM

# Questions?

